



The Neo4j Migration Guide

Table of Contents

1. Important information	3
1.1. Understanding upgrades and migration	3
1.2. Supported paths	4
2. Prepare to migrate	6
3. Surface changes	8
3.1. Security	8
3.2. Changes to other configuration settings	8
3.3. Removal of REST API	9
3.4. HTTP API endpoints	9
3.5. Cypher syntax	9
3.6. Database naming rules	10
3.7. Procedures	10
3.8. Authentication and authorization	12
3.9. Logs	13
3.10. Metrics	14
3.11. Cluster discovery	14
3.12. Cluster REST endpoints	14
3.13. JMX	15
3.14. Index migration	15
3.15. Tools	15
3.16. Backups	15
3.17. Embedded layout	16
3.18. Core Java API	16
4. Migrate a single instance (offline)	20
5. Migrate a Causal Cluster (offline)	22
6. Migrate Neo4j drivers	23
6.1. New driver releases	23
6.2. Compatibility	23
6.3. What's new?	24
6.4. Breaking changes	25
6.5. Back-pressure	30
6.6. Multiple databases	31
6.7. Configure SSL Policy for Bolt server and HTTPS server	32
6.8. Additional URI Schemes	33
7. Tutorial: Back up and restore a database	35
7.1. Prepare to back up the database	35
7.2. Back up the database	36
7.3. Restore the database backup on a 4.x standalone instance	37
7.4. Recreate the database users and roles	38
8. Tutorial: Back up and restore a database in a Causal Cluster	39
8.1. Prepare to back up the database	39
8.2. Back up the database	41
8.3. Migrate the database backup to a 4.x standalone instance	41
8.4. Back up the migrated database	42
8.5. Restore the migrated database on all cluster members	42
8.6. Create the database backup on the cluster leader	43

8.7. Recreate the database users and roles	43
9. Tutorial: Back up and copy a database	45
9.1. Check your 3.5 database store usage	45
9.2. Back up your 3.5 database	46
9.3. Copy your 3.5 database backup to 4.x Neo4j	46
9.4. Create your compacted backup	48
Appendix A: Classes removed from public API.	51
Appendix B: External dependencies	61

This guide describes how to migrate from Neo4j version 3.5 to 4.0.



Migration can be done from Neo4j version 3.5.latest to 4.0.latest.

If a later version is desired, follow the upgrade steps in the [Upgrade](#) chapter of the Operations Manual, once the migration is complete.

This guide describes the following:

- [Important information](#) — Information such as [supported migration paths](#) and [known limitations](#).
- [Prepare to migrate](#) — A checklist for migrating from version 3.5 to 4.x.
- [Surface changes](#) — A list of breaking changes to the Neo4j surface between version 3.5 and 4.x. These include:
 - [Security](#)
 - [Changes to configuration settings](#)
 - [Removal of REST API](#)
 - [HTTP API endpoints](#)
 - [Cypher syntax](#)
 - [Database naming rules](#)
 - [Procedures](#)
 - [Authentication and authorization](#)
 - [Logs](#)
 - [Metrics](#)
 - [Cluster discovery](#)
 - [REST endpoints](#)
 - [JMX](#)
 - [Index migration](#)
 - [Tools](#)
 - [Core Java API](#)
- [Migrate a single instance \(offline\)](#) — Instructions for how to migrate a single instance from Neo4j version 3.5 to 4.0.
- [Migrate a Causal Cluster \(offline\)](#) — Instructions for how to migrate a Causal Cluster from Neo4j version 3.5 to 4.0.
- [Tutorial: Back up and restore a database \(online\)](#) — A detailed example of an online backup and restore of a database.
- [Tutorial: Back up and restore a database in a Causal Cluster](#) — A detailed example of an online backup and restore of a database in a Causal Cluster.
- [Tutorial: Back up and copy a database \(online\)](#) — A detailed example of an online backup and copy of a database.
- [Migrate Neo4j drivers](#) — Information for migrating Neo4j drivers from version 1.7 to 4.0.

Appendix:

- [Classes removed or excluded from the public API](#) — A list of the classes that have been removed or excluded from the public API between Neo4j version 3.5 and 4.0.
- [External dependencies](#) — A list of the external dependencies in Neo4j version 4.0.

Who should read this?

This migration guide is written for:

- the engineer performing the Neo4j production migration.
- the operations engineer supporting and maintaining the Neo4j production database.
- the enterprise architect researching database migration.
- the infrastructure architect planning the Neo4j production migration.
- the enterprise data security manager responsible for the company's strategy for role-based access control.

Chapter 1. Important information

This section explains the difference between upgrade and migration, as well as the supported paths for keeping your deployment up-to-date.

Keeping your Neo4j deployment always up-to-date ensures that you are provided with the latest improvements in performance, security, and bug fixes.

1.1. Understanding upgrades and migration

Before you start migrating to a newer version of Neo4j, you should have a clear understanding of upgrades and migration.

1.1.1. Upgrade

The process of upgrading an existing Neo4j deployment (single or clustered) to a newer **minor or release version** of Neo4j, when such process **does not require** changes to the configuration or to the applications that use Neo4j. For example:

- Between adjacent patch releases within the same major and minor version, e.g., 3.5.0 to 3.5.1, 4.2.1 to 4.2.2.
- Between non-adjacent patch releases within the same major and minor version, e.g., 3.5.0 to 3.5.10, 4.2.1 to 4.2.3.
- Between adjacent minor versions, e.g., 4.0.0 to 4.1.3.

What needs to be upgraded

- The Neo4j product.
- The store formats (the file formats on disk, which usually auto-upgrade at server startup, but sometimes an older version may remain on a running server).
- The **system** database schema, which restructures the contents of that database. (From 4.1 onwards) This is done automatically on standalone server and in offline cluster upgrades (where you temporarily set each server as standalone), but need to be run manually on rolling upgrades after the other two upgrades (product and stores) finish.

1.1.2. Migration

The process of migrating an existing Neo4j deployment (single or clustered) to a newer **major** Neo4j version, when such process **requires** a review of the configuration(s) and the applications that use Neo4j. For example:

- Between major versions, e.g., from 3.5.13 to 4.0.10.

What needs to be migrated

- The Neo4j product.
- The store formats (the file formats on disk, which usually auto-upgrade at server startup, but sometimes an older version may remain on a running server).
- The **system** database schema, which restructures the contents of that database.
- Configuration settings - the configuration changes from the newer version must be applied to the old configuration file.

- Application code - the source code of your application(s) must be updated as per the new version in order to work properly.
- Custom plugins - all custom plugins must be compatible with the new version of Neo4j.

Limitations

- Neo4j does not support downgrades.
- A Neo4j migration must be performed as an isolated operation.
- Requires Neo4j DBMS downtime.

1.1.3. Store copy

You can use the `neo4j-admin copy` command to migrate from Neo4j 3.5.x directly to 4.2.y, skipping the intermediate steps of 3.5.latest \square 4.0.latest \square 4.1.latest \square 4.2.latest. Those steps are needed to migrate the *schema*, but since the `neo4j-admin copy` command does not copy the schema store at all, they are not needed. However, if there is a schema defined, you have to recreate it by running the commands that the `neo4j-admin copy` operation outputs. The `neo4j-admin copy` command can be applied only to an **offline** database.



Your copied node IDs will be the same, but the relationships will get new IDs. Therefore, if you want to preserve the relationship IDs, follow the sequential path 3.5.any \square 4.0.any \square 4.1.any \square 4.2.any.

1.2. Supported paths

Neo4j supports two paths for migrating from an earlier version of Neo4j to the latest:

- sequential — 3.5.any \square 4.0.any \square 4.1.any \square 4.2.any
- direct — 3.5.any \square 4.x.any

Depending on your current version, backup and restore strategy, and the version you want to go to, you can choose the most appropriate path for you. You can also move from Community to Enterprise Edition (vice versa is not supported), as well as from a single instance to a Causal cluster.

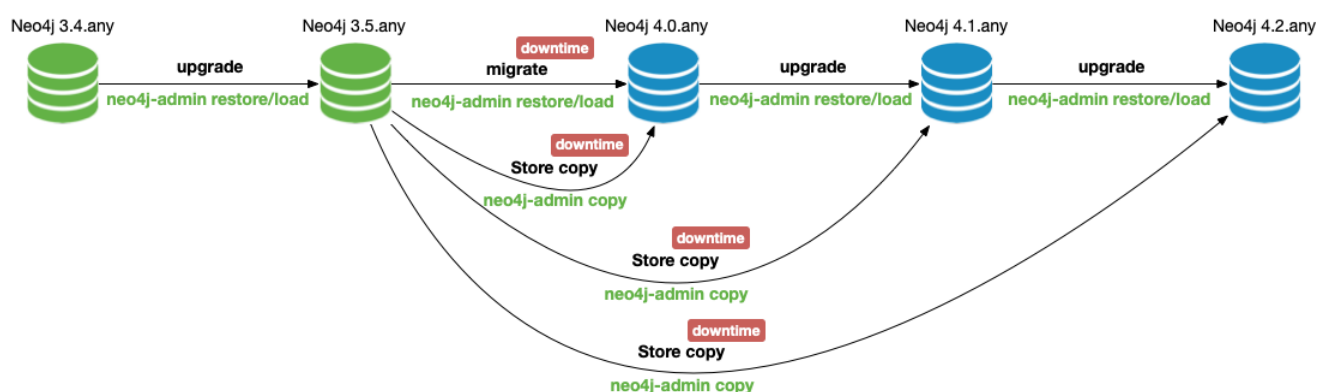


Table 1. Documentation for completing your upgrade/migration

Starting version	Target version	Supported path	Operations	Documentation
3.4.any	3.5.any	3.4.any \square 3.5.any	Upgrade	Operations Manual v3.5 \square Upgrade

Starting version	Target version	Supported path	Operations	Documentation
3.5.any	4.0.any	3.5.any → 4.0.any	Migration	For a standalone instance, see Tutorial: Back up and restore a database . For a Causal cluster, see Tutorial: Back up and restore a database in a Causal Cluster .
3.5.any	4.1.any	3.5.any → 4.0.any → 4.1.any	Migration and upgrade	For a standalone instance, see Tutorial: Back up and restore a database . For a Causal cluster, see Tutorial: Back up and restore a database in a Causal Cluster .
3.5.any	4.2.any	3.5.any → 4.0.any → 4.1.any → 4.2.any	Migration and upgrades	For a standalone instance, see Tutorial: Back up and restore a database . For a Causal cluster, see Tutorial: Back up and restore a database in a Causal Cluster .
3.5.any	4.0.any	3.5.any → 4.0.any	Store copy	Tutorial: Back up and copy a database
3.5.any	4.1.any	3.5.any → 4.1.any	Store copy	Tutorial: Back up and copy a database
3.5.any	4.2.any	3.5.any → 4.2.any	Store copy	Tutorial: Back up and copy a database
4.0.any	4.1.any	4.0.any → 4.1.any	Upgrade	Operations Manual v4.1 → Upgrade
4.0.any	4.2.any	4.0.any → 4.1.any → 4.2.any	Upgrades	Operations Manual v4.1 → Upgrade Operations Manual v4.2 → Upgrade
4.1.any	4.2.any	4.1.any → 4.2.any	Upgrade	Operations Manual v4.2 → Upgrade

Chapter 2. Prepare to migrate

This chapter provides a checklist of things to prepare before performing a migration from Neo4j version 3.5 to 4.0.



The supported upgrade path to Neo4j versions beyond 4.0 is **3.5.latest** → **4.0.latest** → **4.1** → **4.2**. Please note that Neo4j versions 4.0 and beyond, requires Java 11.

Follow this checklist in order to ensure that you are well prepared before you start a production migration from Neo4j version 3.5 to 4.0:

Review Release Notes

To view the details of the changes that are included in each version, see the [Release Notes](#).

Apply configuration changes

Prepare the contents of [neo4j.conf](#) to be used for the migrated database. If you are migrating a Causal Cluster, do this for each of the members in the cluster.

- Review the section [Changes to other configuration settings](#) and update all applicable configuration settings.
- It is also useful to inspect the current configuration file and take note of any non-default settings. When upgrading, it is particularly important to note any custom values of the settings `dbms.directories.data` and `dbms.default_database`. In cluster installations, pay attention to cluster-specific configuration settings, which may be different on different cluster members.



Some configuration settings that have changed names are automatically migrated to the new setting names during startup. When this happens, it is logged in [neo4j.log](#). The automatic migration is not permanent, so if it is not changed in [neo4j.conf](#), it will take place each time at startup. When the deprecated setting name are subsequently removed, unexpected problems may occur. It is therefore strongly recommended to update all relevant configuration settings at the time of the migration to 4.0.

Upgrade application code

Review the changes outlined in this guide and apply the necessary changes to your source code. How much development time is required to update the code will depend on the particular application. Make sure to test the application code thoroughly.

Upgrade custom plugins

Check the [plugins](#) directory to verify whether custom plugins are used in your deployment. Ensure that any plugins are compatible with Neo4j 4.x.

Plan disk space requirements

A migration requires substantial free disk space, as it makes an entire copy of the database. For the migration, make sure to make available an additional 50% * `size_of(database directory)`. In a default configuration, the `database directory` is `databases/neo4j`, which is located in the [data](#) directory. In addition to this, do not forget to reserve the disk space needed for the pre-migration backup.

The migrated database may require slightly larger data files overall.

Perform a test migration

Based on the findings in this chapter, allocate a production-like test environment for the migration and do a test migration. The test migration will give you valuable information about the time

required for the production migration, as well as potential additional action points, such as migration of plugins and application code.

Review the logs

The [neo4j.log](#) file contains valuable information on how many steps the migration will involve and how far it has progressed. For large migrations, it is a good idea to monitor this log continuously. Below is a sample of what the log may look like:

```
2018-09-18 13:24:23.243+0000 INFO Starting...
2018-09-18 13:24:24.262+0000 INFO Initiating metrics...
2018-09-18 13:24:24.488+0000 INFO Starting upgrade of database
2018-09-18 13:24:24.538+0000 INFO Migrating Indexes (1/5):
2018-09-18 13:24:24.542+0000 INFO 10% completed
2018-09-18 13:24:24.543+0000 INFO 20% completed
2018-09-18 13:24:24.543+0000 INFO 30% completed
...
...
...
2018-09-18 13:24:24.574+0000 INFO Migrating Counts store (5/5):
2018-09-18 13:24:24.574+0000 INFO 10% completed
2018-09-18 13:24:24.574+0000 INFO 20% completed
2018-09-18 13:24:24.575+0000 INFO 30% completed
...
...
...
2018-09-18 13:24:24.576+0000 INFO 100% completed
2018-09-18 13:24:24.584+0000 INFO Successfully finished upgrade of database
```

Chapter 3. Surface changes

This chapter describes breaking changes to the Neo4j surface when migrating from Neo4j version 3.5 to 4.0 and beyond.

3.1. Security

In version 4.x, the SSL framework has been reworked and the legacy SSL system has been deprecated. The following table outlines the security changes:

v3.x	v4.x
The configuration settings <code>dbms.security.property_level.enabled</code> and <code>dbms.security.property_level.blacklist</code> are used to disallow properties.	These configuration settings have been discontinued. The blocking functionality is replaced by the Cypher <code>DENY</code> command. Note that the <code>DENY</code> command must be applied while Neo4j is running. For details, see Cypher Manual ▯ Security ▯ Graph and sub-graph access control .
<code>dbms.security.auth_provider</code>	This setting is replaced by two new settings: <code>dbms.security.authentication_providers</code> and <code>dbms.security.authorization_providers</code> .
<code>dbms.connector.https.enabled</code> is set to <code>true</code> by default.	This setting is no longer <code>true</code> by default. To enable Neo4j to listen for incoming connections on the HTTPS port, you have to configure this setting to <code>true</code> .
The different communication channels are secured independently from each other, using the following configuration settings: <code>bolt.ssl_policy=<policy name></code> <code>https.ssl_policy=<policy name></code> <code>causal_clustering.ssl_policy=<policy name></code> <code>dbms.backup.ssl_policy=<policy name></code>	These settings have been replaced by the setting <code>dbms.ssl_policy.<scope>.enabled=true</code> , where <code><scope></code> substitutes the communication channel (<code>bolt</code> , <code>https</code> , <code>cluster</code> , and <code>backup</code>).
SSL support for Bolt and HTTPS using the legacy SSL system. The <code>dbms.directories.certificates</code> setting is used to explicitly configure the directory that stores the private key and certificate files.	The legacy SSL system has been deprecated and the <code>dbms.directories.certificates</code> setting has been removed. It is recommended to use the standard SSL configuration.
<code>dbms.ssl_policy.*.allow_key_generation</code>	This setting has been removed. Neo4j no longer automatically generates a self-signed certificate.



For further details on the SSL framework changes, see [Operations Manual ▯ SSL framework](#)

3.2. Changes to other configuration settings

Previous name	Change	New name (if applicable)
<code>dbms.active_database</code>	Renamed	<code>dbms.default_database</code>
<code>dbms.connectors.default_listen_address</code>	Renamed	<code>dbms.default_listen_address</code>
<code>dbms.connectors.default_advertised_address</code>	Renamed	<code>dbms.default_advertised_address</code>
<code>dbms.backup.address</code>	Renamed	<code>dbms.backup.listen_address</code>
<code>dbms.logs.query.enabled</code>	This is no longer a boolean setting. Valid values are: <code>OFF</code> , <code>INFO</code> or <code>VERBOSE</code> .	

Previous name	Change	New name (if applicable)
<code>causal_clustering.cluster_routing_ttl</code>	Renamed	<code>dbms.routing_ttl</code>
<code>causal_clustering.middleware_logging.level</code>	This setting has been renamed, and valid values are: <code>DEBUG</code> , <code>INFO</code> , <code>WARN</code> , <code>ERROR</code> or <code>NONE</code>	<code>causal_clustering.middleware.logging.level</code>
<code>causal_clustering.disable_middleware_logging</code>	This setting is removed. Set <code>causal_clustering.middleware_logging.level=OFF</code> to disable middleware logging.	
<code>metrics.neo4j.logrotation.enabled</code>	Renamed	<code>metrics.neo4j.logs.enabled</code>
<code>metrics.enabled</code>	This setting no longer changes the default values of the individual metrics. Instead it turns off the whole metrics module.	

3.3. Removal of REST API

The REST API has been removed in Neo4j 4.0. Cypher and procedures should be used instead, either via the HTTP API, or via Bolt using the official drivers.

The following HTTP endpoints were deprecated in Neo4j 3.4 and have now been removed:

HTTP endpoints
<code>/db/data/branch</code>
<code>/db/data/cypher</code>
<code>/db/data/index/node</code>
<code>/db/data/index/relationship</code>
<code>/db/data/labels</code>
<code>/db/data/node</code>
<code>/db/data/relationship</code>
<code>/db/data/relationship/types</code>
<code>/db/data/schema/constraint</code>
<code>/db/data/schema/index</code>
<code>/db/data/schema/relationship/constraint</code>

3.4. HTTP API endpoints

The HTTP API endpoints have been updated to accommodate multi database features. For example, the URI to begin a transaction has changed from: `http://localhost:7474/db/data/transaction` to: `http://localhost:33471/db/neo4j/tx`.

More generally, the HTTP API endpoints follow the pattern: `http://localhost:33471/db/<database_name>/tx`.

3.5. Cypher syntax

- All changes in the Cypher language syntax are detailed in [Cypher Manual ▯ Removals, deprecations, additions and extensions](#). Please review it thoroughly and make necessary changes in your code.
- We would like to draw some extra attention to the fact that the parameter syntax `{parameter}` is

completely removed and has been replaced by the syntax `$parameter`.

3.6. Database naming rules

With the introduction of multiple databases, the rules for naming a database have changed. For example, it is no longer possible to use an underscore in a database name. For a full list of naming rules, please see [Operations Manual](#) ▢ [Administrative commands](#).

3.7. Procedures

The following procedures have been refactored:

Old procedure	New procedure	Comment
<code>db.awaitIndex (indexId :: INTEGER?, timeOutSeconds = 300 :: INTEGER?) :: VOID</code>	<code>db.awaitIndex (indexName :: STRING?, timeOutSeconds = 300 :: INTEGER?) :: VOID</code>	Indexes are now uniquely identified by name, instead of ID.
<code>dbms.cluster.overview() :: (id :: STRING?, addresses :: LIST? OF STRING?, role :: STRING?, groups :: LIST? OF STRING?, database :: STRING?)</code>	<code>dbms.cluster.overview() :: (id :: STRING?, addresses :: LIST? OF STRING?, databases :: MAP?, groups :: LIST? OF STRING?)</code>	Shows roles for all databases.
<code>dbms.cluster.role() :: (role :: STRING?)</code>	<code>dbms.cluster.role (database :: STRING?) :: (role :: STRING?)</code>	Takes <code>database</code> name as parameter.
<code>dbms.cluster.routing.getRoutingTable (context :: MAP?) :: (ttl :: INTEGER?, servers :: LIST? OF MAP?)</code>	<code>dbms.cluster.routing.getRoutingTable (context :: MAP?, database = null :: STRING?) :: (ttl :: INTEGER?, servers :: LIST? OF MAP?)</code>	Takes <code>database</code> name as parameter.
<code>db.createIndex (index :: STRING?, providerName :: STRING?) :: (index :: STRING?, providerName :: STRING?, status :: STRING?)</code>	<code>db.createIndex (indexName :: STRING?, labels :: LIST? OF STRING?, properties :: LIST? OF STRING?, providerName :: STRING?, config = {} :: MAP?) :: (name :: STRING?, labels :: LIST? OF STRING?, properties :: LIST? OF STRING?, providerName :: STRING?, status :: STRING?)</code>	Used to take the index pattern <code>":Label(prop)"</code> as an argument, and now takes <code>labels</code> and <code>properties</code> as separate lists. Those are also yielded as result. Now needs to be given an <code>indexName</code> . Can now take index settings as a map. This is optional.
<code>db.createUniquePropertyConstraint (index :: STRING?, providerName :: STRING?) :: (index :: STRING?, providerName :: STRING?, status :: STRING?)</code>	<code>db.createUniquePropertyConstraint (constraintName :: STRING?, labels :: LIST? OF STRING?, properties :: LIST? OF STRING?, providerName :: STRING?, config = {} :: MAP?) :: (name :: STRING?, labels :: LIST? OF STRING?, properties :: LIST? OF STRING?, providerName :: STRING?, status :: STRING?)</code>	Used to take the index pattern <code>":Label(prop)"</code> as an argument, and now takes <code>labels</code> and <code>properties</code> as separate lists. Those are also yielded as result. Now needs to be given a <code>constraintName</code> . Can now take index settings as a map. This is optional.

Old procedure	New procedure	Comment
<pre>db.createNodeKey (index :: STRING?, providerName :: STRING?) :: (index :: STRING?, providerName :: STRING?, status :: STRING?)</pre>	<pre>db.createNodeKey (constraintName :: STRING?, labels :: LIST? OF STRING?, properties :: LIST? OF STRING?, providerName :: STRING?, config = {} :: MAP?) :: (name :: STRING?, labels :: LIST? OF STRING?, properties :: LIST? OF STRING?, providerName :: STRING?, status :: STRING?)</pre>	<p>Used to take the index pattern <code>":Label(prop)"</code> as an argument, and now takes <code>labels</code> and <code>properties</code> as separate lists. Those are also yielded as result.</p> <p>Now need to be given a <code>constraintName</code>.</p> <p>Can now take index settings as a map. This is optional.</p>
<pre>db.indexes() :: (description :: STRING?, indexName :: STRING?, tokenNames :: LIST? OF STRING?, properties :: LIST? OF STRING?, state :: STRING?, type :: STRING?, progress :: FLOAT?, provider :: MAP?, id :: INTEGER?, failureMessage :: STRING?)</pre>	<pre>db.indexes() :: (id :: INTEGER?, name :: STRING?, state :: STRING?, populationPercent :: FLOAT?, uniqueness :: STRING?, type :: STRING?, entityType :: STRING?, labelsOrTypes :: LIST? OF STRING?, properties :: LIST? OF STRING?, provider :: STRING?)</pre>	<p>Rename <code>indexName</code> to <code>name</code>.</p> <p>Rename <code>tokenNames</code> to <code>labelsOrTypes</code>.</p> <p>Rename <code>progress</code> to <code>populationPercent</code>.</p> <p>Field <code>type</code> used to describe entity type (node or relationship), uniqueness, and index type. This splits up into <code>type</code>, <code>uniqueness</code>, and <code>entityType</code>.</p> <p>Field <code>provider</code> is now a string instead of a map.</p> <p>Removed <code>description</code> in favor of <code>db.schemaStatements</code>.</p> <p>Moved <code>failureMessage</code> to procedure <code>db.indexDetails</code>.</p>
<pre>db.resampleIndex (index :: STRING?) :: VOID</pre>	<pre>db.resampleIndex (indexName :: STRING?) :: VOID</pre>	<p>Indexes are now uniquely identified by name, instead of index pattern <code>":Label(prop)"</code>.</p>

The following are new procedures:

New procedure	Comment
<pre>db.indexDetails (indexName :: STRING?) :: (id :: INTEGER?, name :: STRING?, state :: STRING?, populationPercent :: FLOAT?, uniqueness :: STRING?, type :: STRING?, entityType :: STRING?, labelsOrTypes :: LIST? OF STRING?, properties :: LIST? OF STRING?, provider :: STRING?, indexConfig :: MAP?, failureMessage :: STRING?)</pre>	<p>For the specified index all information included by <code>db.indexes</code> together with <code>indexConfig</code> and <code>failureMessage</code>.</p>
<pre>db.schemaStatements () :: (name :: STRING?, type :: STRING?, createStatement :: STRING?, dropStatement :: STRING?)</pre>	<p>Get all create and drop statements needed to exactly replicate the schema rules (indexes and constraints) for this database.</p>

New procedure	Comment
<code>db.ping()</code>	This procedure can be used by client side tooling to test whether they are correctly connected to a database. The procedure is available in all databases and always returns true. A faulty connection can be detected by not being able to call this procedure.

3.8. Authentication and authorization

3.8.1. Deprecated and removed security procedures

In 3.x, authentication and authorization was managed via the built-in `dbms.security` procedures. In 4.x, these procedures still exist but are deprecated. If you still want to use them, they must now be run in a session towards the `system` database, and cannot be followed by `YIELD`. There are two options for rewriting your code and routines for managing authentication and authorization. The first of these is recommended:

1. Rewrite the procedures to the corresponding Cypher administration commands, using the [the conversion guide below](#).
2. Run the procedures in a session towards the `system` database and replace any `YIELD` parts by post-processing on the application side.



The procedure `dbms.security.changePassword(password, requirePasswordChange)` has been entirely removed since the corresponding Cypher administration command also requires the old password, and thus is more secure.

The following table is a conversion guide between the security procedures and the Cypher administration commands. For more info about the administration commands, see [Cypher Manual](#) [User and role management](#).

Procedure	Administration command
<code>dbms.security.createUser</code>	<code>CREATE USER</code>
<code>dbms.security.deleteUser</code>	<code>DROP USER</code>
<code>dbms.security.changePassword</code>	<code>ALTER CURRENT USER SET PASSWORD</code>
<code>dbms.security.listUsers</code>	<code>SHOW USERS</code>
<code>dbms.security.changeUserPassword</code>	<code>ALTER USER</code>
<code>dbms.security.suspendUser</code>	<code>ALTER USER</code>
<code>dbms.security.activateUser</code>	<code>ALTER USER</code>
<code>dbms.security.addRoleToUser</code>	<code>GRANT ROLE TO USER</code>
<code>dbms.security.removeRoleFromUser</code>	<code>REVOKE ROLE FROM USER</code>
<code>dbms.security.listRoles</code>	<code>SHOW ROLES</code>
<code>dbms.security.listRolesForUser</code>	<code>SHOW USERS</code>
<code>dbms.security.listUsersForRole</code>	<code>SHOW ROLES WITH USERS</code>
<code>dbms.security.createRole</code>	<code>CREATE ROLE</code>
<code>dbms.security.deleteRole</code>	<code>DROP ROLE</code>

3.8.2. Removal of flat files for authentication and authorization

In 3.x, authentication and authorization were managed in flat files. Users in the *auth* file and roles and role assignments in the *roles* file will be automatically migrated to the **system** database when upgrading from Neo4j 3.5 to Neo4j 4.0.

The Neo4j admin commands **set-initial-password** and **set-default-admin** continue to work in 4.0 and write to the same files as in 3.x. Any content in these files will be considered on the first start of Neo4j after upgrading from 3.5. You can run these commands before upgrading the Neo4j installation, or after, as long as they are run before completing the migration of the database files which is done at first start of the new installation.



The command **set-initial-password** will only be applied if the default user **neo4j** with the default password is the only user present, while **set-default-admin** will only be applied when no roles are present.

The use of *auth* and *role* files in Neo4j 3.x meant that multiple databases could have different user and role configurations. In addition, a single database configured in a cluster could have different *auth* and *role* settings on each instance of the cluster. Neo4j 4.0 allows multiple databases to run within a single instance, or in a cluster. If you are bringing multiple databases together from multiple Neo4j 3.5 installations, or if you are upgrading a cluster with multiple instances, you need to manually merge the *auth* and *role* files before the migration.

It is still possible to have different security configurations per database after the migration, but this needs to be managed through the granting of privileges and roles specific to databases after the migration. The built-in roles from 3.5 still exist, but will apply to all databases after the migration, unless explicitly modified using the new security administration commands. The ability to manage database specific roles and privileges is described in more detail in [Cypher Manual ▯ Administration](#).

It is no longer possible to have different security privileges on different instances of a cluster. The entire cluster shares the privileges configured in the **system** database using Cypher administration commands. In practice this means that users have the same privileges regardless of which server in a cluster they access.

3.9. Logs

Relevant logs produced by Neo4j will now have a prefix which indicates the database to which the log line pertains. Such log lines will have the database name printed prior to the regular text. For example, **[neo4j]** or **[system]**.

Example 1. Some log lines for the system database

```
2019-12-02 22:27:41.820+0000 INFO [o.n.k.d.Database] [system] No check point found in transaction log
2019-12-02 22:27:41.820+0000 INFO [o.n.k.d.Database] [system] Recovery required from position
LogPosition{logVersion=0, byteOffset=64}
2019-12-02 22:27:41.820+0000 INFO [o.n.k.r.Recovery] [system] 10% completed
2019-12-02 22:27:41.820+0000 INFO [o.n.k.r.Recovery] [system] 20% completed
2019-12-02 22:27:41.820+0000 INFO [o.n.k.r.Recovery] [system] 30% completed
...
```

Other log lines might relate to the DBMS as a whole, or be logged by a component that lives on a higher level but still operates on a particular database. For example:

Example 2. Some log lines from the Core database manager starting the Neo4j database.

```
2019-12-02 22:27:41.964+0000 INFO [c.n.c.c.CoreDatabaseManager] Creating 'neo4j' database.
2019-12-02 22:27:41.967+0000 INFO [c.n.c.c.CoreDatabaseManager] Starting 'neo4j' database.
...
```

3.10. Metrics

In 4.x, there are two types of metrics: global metrics and database-local metrics. The metric naming is different in 4.x compared to 3.x. For details about available metrics and the new naming patterns, please refer to [Operations Manual ▯ Metrics](#).



Please note that in 4.2 the metrics that are enabled by default have been changed.

Any specific metrics that you want to be enabled **must** be specified in the `metrics.filter`.

Additionally in the 4.2 release, metrics are no longer exposed via JMX by default. These can be enabled by adding `metrics.jmx.enabled=true` to `neo4j.conf`.

For more information, see [Operations Manual ▯ Enable metrics logging](#)

3.11. Cluster discovery

Cluster discovery is now implemented on top of Akka, instead of Hazelcast, and a few minor changes have been made as part of this transition:

- The `discovery_advertised_address` hostname and port must exactly match those configured for the discovery of other members.

When `discovery_type=LIST` is used, this means that it is the list of addresses in `initial_discovery_members` which must match the respective advertised addresses of each server.

When using any other discovery types (DNS, SRV, K8S), then it is the configuration in the external service which must match.



Please note that by default your `discovery_advertised_address` is a combination of the default port assigned to that config, and the hostname assigned to `default_advertised_address`.

- Connections are now opened from Cores to Read Replicas, in addition to vice versa, so therefore the advertised discovery port must be **open** on Read Replicas.

3.12. Cluster REST endpoints

The REST endpoints have moved and now exist per database:

Old endpoint	New endpoint
<code>/db/manage/server/causalclustering/writable</code>	<code>/db/<databasename>/cluster/writable</code>
<code>/db/manage/server/causalclustering/read-only</code>	<code>/db/<databasename>/cluster/read-only</code>
<code>/db/manage/server/causalclustering/available</code>	<code>/db/<databasename>/cluster/available</code>
<code>/db/manage/server/causalclustering/status</code>	<code>/db/<databasename>/cluster/status</code>

3.13. JMX

In 3.x, Neo4j exposed several JMX MBeans in order to provide some monitoring information in addition to the metrics exposed by Neo4j. In some instances, the provided data was incomplete or incorrect, and in some cases different beans even provided conflicting information. All of the previous JMX endpoints (`org.neo4j:*`) have been removed and are replaced by a new set of beans (`neo4j.metrics:*`) that expose exactly the same information as the corresponding Neo4j metrics.

JMX MBeans are available only in Enterprise Edition.

3.14. Index migration

- Indexes are automatically migrated to the most recent index provider during migration.

Depending on what index providers were used previously, the migration of indexes may change the distribution of memory utilization. In a database with many indexes, a significant amount of memory may have been reserved for Lucene. After the migration, it could be necessary to allocate some of that memory to the page cache instead. For a detailed description on how memory is allocated and used, refer to [Operations Manual ▯ Memory configuration](#). Use `neo4j-admin memrec --database` to inspect the database before and after migration.

Changes have been made to how large a key can be in a b-tree index. These changes are only relevant for indexes that use index provider `lucene-1.0` or `lucene+native-1.0` in 3.5, and hold large strings or large arrays. For a detailed description of this change, please refer to [Operations Manual ▯ Index migration](#).

- Support for explicit indexes has been removed and the functionality has been replaced by full-text indexes. For details, see [Cypher Manual ▯ Indexes to support full-text search](#).

3.15. Tools

Database specific commands provided by `neo4j-admin` now support `--database`, which can be used to specify a database for a specified operation.

In cases when the `--database` option is not specified, `neo4j` will be used as the default database.

Also, there is a slight syntax change when adding options to `neo4j-admin import`. To add a label or relationship type to all nodes or relationships in an import file, the syntax is: `neo4j-admin import --nodes=[<label>[:<label>]...]=<files>...` and `neo4j-admin import --relationships=[<type>]=<files>...`.

In addition, with the introduction of multiple databases, it is important to remember that if importing to a new database, it has to be explicitly created before the imported data can be accessed. For example, if the database is called `importeddb`, after data has been imported to it, create it with the following query:

```
:use system
CREATE DATABASE importeddb
```

For naming rules, please see [Cypher Manual ▯ Naming rules and recommendations](#) and [Operations Manual ▯ Administrative commands](#).

3.16. Backups

Backups must now be taken of all databases.

A default installation has two databases, named `system` and `neo4j` respectively. Use the `--database` option of the `neo4j-admin backup` command to specify the database to backup. For more information, see [Operations Manual](#) ▢ [Back up an online database](#).

The `--name` parameter has been removed. It was previously used to specify the last part of the path when using `--backup-dir`. The last part of the path is now inferred from the `--database` parameter, which is used to specify the database name on the server. You are therefore no longer able to specify the last part of the path.

If you previously used `--name` for customizing the backup path, for example by including a timestamp, then an alternative is to now use `--backup-dir` instead.

3.17. Embedded layout

To support multiple databases in embedded, the store files, transaction files and log files no longer reside in the base directory. Instead, files are separated per database in separate directories.

3.18. Core Java API

3.18.1. JDK 11

Neo4j 4.0 is the first major release that requires JDK 11. Custom extensions and procedures can also be compiled now for JDK 11 (for example `-target 11`). It is generally recommended to use the latest available JDK 11 in order to access available fixes and leverage performance improvements.

3.18.2. Classes removed or excluded from the public API

Please refer to [Classes removed from public API](#) for a complete list of classes removed or excluded from the public API.

3.18.3. Renamed classes

The following classes have been renamed:

Old class name	New class name
<code>org.neo4j.graphdb.factory.GraphDatabaseSettings.BoltConnector</code>	<code>org.neo4j.configuration.connectors.BoltConnector</code>
<code>org.neo4j.graphdb.factory.GraphDatabaseSettings.BoltConnector.EncryptionLevel</code>	<code>org.neo4j.configuration.connectors.BoltConnector.EncryptionLevel</code>
<code>org.neo4j.kernel.configuration.HttpConnector</code>	<code>org.neo4j.configuration.connectors.HttpConnector</code>
<code>org.neo4j.graphdb.factory.GraphDatabaseSettings</code>	<code>org.neo4j.configuration.GraphDatabaseSettings</code>
<code>org.neo4j.graphdb.factory.GraphDatabaseSettings.SchemaIndex</code>	<code>org.neo4j.configuration.GraphDatabaseSettings.SchemaIndex</code>
<code>org.neo4j.backup.OnlineBackup</code>	<code>com.neo4j.backup.OnlineBackup</code>
<code>org.neo4j.helpers.SocketAddress</code>	<code>org.neo4j.configuration.helpers.SocketAddress</code>
<code>org.neo4j.graphdb.event.TransactionEventHandler</code>	<code>org.neo4j.graphdb.event.TransactionEventListener</code>
<code>org.neo4j.graphdb.factory.GraphDatabaseSettings.TransactionStateMemoryAllocation</code>	<code>org.neo4j.configuration.GraphDatabaseSettings.TransactionStateMemoryAllocation</code>
<code>org.neo4j.graphdb.index.fulltext.AnalyzerProvider</code>	<code>org.neo4j.graphdb.schema.AnalyzerProvider</code>
<code>org.neo4j.server.security.enterprise.auth.plugin.api.AuthProviderOperations</code>	<code>com.neo4j.server.security.enterprise.auth.plugin.api.AuthProviderOperations</code>

Old class name	New class name
<code>org.neo4j.server.security.enterprise.auth.plugin.api.AuthToken</code>	<code>com.neo4j.server.security.enterprise.auth.plugin.api.AuthToken</code>
<code>org.neo4j.server.security.enterprise.auth.plugin.api.AuthenticationException</code>	<code>com.neo4j.server.security.enterprise.auth.plugin.api.AuthenticationException</code>
<code>org.neo4j.server.security.enterprise.auth.plugin.api.P</code> <code>redefinedRoles</code>	<code>com.neo4j.server.security.enterprise.auth.plugin.api.P</code> <code>redefinedRoles</code>
<code>org.neo4j.server.security.enterprise.auth.plugin.spi.AuthInfo</code>	<code>com.neo4j.server.security.enterprise.auth.plugin.spi.AuthInfo</code>
<code>org.neo4j.server.security.enterprise.auth.plugin.spi.AuthPlugin</code>	<code>com.neo4j.server.security.enterprise.auth.plugin.spi.AuthPlugin</code>

3.18.4. Changes to the API

org.neo4j.graphdb.schema

Neo4j 4.0 comes with significant changes in schema and indexes. Most of the related classes have additional possibilities. Changes include:

- Starting with 4.0, all of the indexes are named. The name of an index can be retrieved using `getName()` call on `IndexDefinition` and `ConstraintDefinition`.
- The definition of an index can be looked up by name using `Schema`.
- Single label and relationship type accessors `getLabel()` and `getRelationshipType()` have been removed from `IndexDefinition`.

Affected classes:

- `org.neo4j.graphdb.schema.ConstraintCreator`
- `org.neo4j.graphdb.schema.ConstraintDefinition`
- `org.neo4j.graphdb.schema.IndexCreator`
- `org.neo4j.graphdb.schema.IndexDefinition`
- `org.neo4j.graphdb.schema.Schema`

org.neo4j.graphdb.event

Transaction event listeners have an updated behavior. Changes include:

- As part of the callback, you will always receive the owning `GraphDatabaseService` as one of the parameters.
- The `beforeCommit` listener method has access to an ongoing transaction over the `transaction` call parameter.
- `DatabaseEventListener` is a new type of listener that has been introduced. Since Neo4j now supports multiple databases you might want to be able to listen to database events from several databases. It can be registered and de-registered in `DatabaseManagementService`.

Affected classes:

- `org.neo4j.graphdb.event.TransactionEventListener`
- `org.neo4j.graphdb.event.DatabaseEventContext`
- `org.neo4j.graphdb.event.DatabaseEventListener`

org.neo4j.helpers

Most of the helpers are no longer part of the public API. The `SocketAddress` helper has minor API changes.

Affected classes:

- `org.neo4j.configuration.helpers.SocketAddress`

com.neo4j.backup

The backup facade has been simplified and adapted to a multi-database environment.

Affected classes:

- `com.neo4j.backup.OnlineBackup`

org.neo4j.configuration

Configuration API has been updated to be typed. It is no longer safe to assume that the configuration is a set of random key-value pairs. All pairs unknown to Neo4j will be rejected. Additionally, some settings have been renamed as well. Please check settings names migration in the corresponding migration manual section.

Affected classes:

- `org.neo4j.configuration.GraphDatabaseSettings`
- `org.neo4j.graphdb.config.Setting`
- `org.neo4j.configuration.connectors.BoltConnector`

org.neo4j.graphdb.Transaction

Transaction API changes are one of the biggest API updates that are part of 4.0. All of the methods that should be executed in transaction have been moved from `GraphDatabaseService` to `Transaction`. This means that if you need to create entities, or access them, you should now be able to find all of the methods in `Transaction`. Additionally, starting with 4.0, transactions are no longer thread-bound. This means that any call to `GraphDatabaseService::beginTransaction()` will create a new independent transaction, even if it called from one thread.

Affected classes:

- `org.neo4j.graphdb.Transaction`

org.neo4j.graphdb.Entity

Starting with 4.0, the `PropertyContainer` interface is removed, and all property-related methods moved to `Entity`. Access to entities should always be transactional. This also means that an entity can only be safely accessed from a transaction where it was created or retrieved.

Affected classes:

- `org.neo4j.graphdb.Entity`
- `org.neo4j.graphdb.Node`
- `org.neo4j.graphdb.Relationship`

org.neo4j.graphdb.GraphDatabaseService

Starting with 4.0, all methods that require transactions are moved to `Transaction`. In addition, a set of `executeTransactionally` methods have been added to provide a convenient way of query executions in a separate transaction.

Affected classes:

- `org.neo4j.graphdb.GraphDatabaseService`

org.neo4j.harness and com.neo4j.harness

Support has been added for official testing support classes. Starting with 4.0, Neo4j provides a set of Junit 4 rules and Junit 5 extensions for community and enterprise users.

Affected classes:

- `com.neo4j.harness.junit.extension.EnterpriseNeo4jExtension`
- `com.neo4j.harness.junit.rule.EnterpriseNeo4jRule`
- `org.neo4j.harness.junit.extension.Neo4j`
- `org.neo4j.harness.junit.extension.Neo4jExtension`
- `org.neo4j.harness.junit.extension.Neo4jExtensionBuilder`
- `org.neo4j.harness.junit.rule.Neo4jRule`

org.neo4j.dbms.api

The top-level Neo4j API has been updated. The main access point that should be used to access individual databases, or perform any database management operations, is called `DatabaseManagementService`. It can be constructed by the Community or Enterprise version of `DatabaseManagementServiceBuilder`.

Example 3. Using DatabaseManagementService

In this example, we are constructing a new `managementService` and a lookup `GraphDatabaseService` for the database named `neo4j`:

```
var managementService = new DatabaseManagementServiceBuilder( homeDirectory ).build();
var databaseService = managementService.database( "neo4j" );
```

Affected classes:

- `org.neo4j.dbms.api.DatabaseManagementService`
- `org.neo4j.dbms.api.DatabaseManagementServiceBuilder`

Chapter 4. Migrate a single instance (offline)

This chapter describes the necessary steps to migrate a single instance from Neo4j version 3.5 to 4.0.



To migrate from 3.5.latest to a version beyond 4.0, the single instance must first be migrated to 4.0 and thereafter upgraded to desired version.

Pre-migration steps

- Refer to [Supported paths](#) regarding supported migration paths.
- Read [Prepare to migrate](#) thoroughly and perform all the steps listed there.

Shutdown and backup

1. If the database is running, shut it down cleanly.
2. Perform and verify backups:
 - Back up [neo4j.conf](#).
 - Back up all the files used for encryption, i.e. private key, public certificate, and the contents of the *trusted* and *revoked* directories. The locations of these are described in [Operations Manual](#) □ [SSL framework](#).
 - Verify that you have a full backup that is stored in a safe location, either using the [online backup tool](#) or [offline backups](#).

Upgrade

1. Install Neo4j 4.0 using one of the following methods, specific to your technology:
 - a. If using a tarball or zipfile for installation:
 - i. Untar or unzip Neo4j 4.0.
 - ii. Transfer the new [neo4j.conf](#) that you prepared in the *Apply configuration changes* step in [Prepare to migrate](#).
 - iii. Set `dbms.allow_upgrade=true` in [neo4j.conf](#) of the 4.0. installation. Neo4j will fail to start without this configuration.
 - iv. Move the files used for encryption from the old installation to the new one.
 - v. Use *Neo4j Admin* tools, such as `neo4j-admin backup` (for online installation), or `neo4j-admin dump` (for offline installation), to move the [data](#) directory from the old installation to the new one. This step is not applicable if you have `dbms.directories.data` pointing to a directory outside of `NEO4J_HOME`.
 - b. If using a Debian or RPM distribution:
 - i. Set `dbms.allow_upgrade=true` in [neo4j.conf](#).
 - ii. Install Neo4j 4.0.
 - iii. When prompted, review the differences between the [neo4j.conf](#) files of the previous version and Neo4j 4.0.
 - iv. Transfer any custom settings to the 4.0 installation, as noted under the *Apply configuration changes* step in [Prepare to migrate](#).

Make sure to preserve `dbms.allow_upgrade=true`, as set in the instruction above.

Neo4j will fail to start without this configuration.

2. Start up Neo4j 4.0. The database migration will take place during startup.

The [neo4j.log](#) file contains valuable information on how many steps the migration will involve and how far it has progressed. For large migrations, it is a good idea to monitor this log continuously.

Post-migration steps

1. When the migration has finished, `dbms.allow_upgrade=true` should be set to `false` or be removed.
2. Restart the database.
3. It is good practice to make a full backup immediately after the migration.

Chapter 5. Migrate a Causal Cluster (offline)

This chapter describes the necessary steps to migrate a Causal Cluster from Neo4j version 3.5 to 4.0.

The migration of a Causal Cluster from Neo4j version 3.5 to 4.0 requires downtime. Therefore, it is recommended to perform a test migration in a production-like environment to get information on the duration of the downtime.



To migrate from 3.5.latest to a version beyond 4.0, the cluster must first be migrated to 4.0 and thereafter upgraded to desired version. See [Operations Manual ▯ Upgrade \(version 4.1\)](#) and/or [Operations Manual ▯ Upgrade \(version 4.2\)](#) for more information.

All the steps must be completed for each cluster member.

Pre-migration steps

- Verify that you have installed Java 11.
- Refer to [Supported paths](#) regarding supported migration paths.
- Read [Prepare to migrate](#) thoroughly and perform all the steps listed there.

Backup

Perform and verify backups:

- Back up [neo4j.conf](#).
- Back up all the files used for encryption, i.e. private key, public certificate, and the contents of the *trusted* and *revoked* directories. The locations of these are described in [Operations Manual ▯ SSL framework](#).
- Verify that you have a full backup that is stored in a safe location, either using the [online backup tool](#) or [offline backups](#).

Migrate

The migration of a cluster follows the same procedure as the cluster upgrade. For detailed instructions, see [Operations Manual ▯ Offline upgrade](#).



For a tutorial on backing up a single database in a running 3.5.x Neo4j instance and restoring it in a running Neo4j 4.0 instance, see [Tutorial: Back up and restore a database](#).

Chapter 6. Migrate Neo4j drivers

This chapter describes the necessary information to migrate Neo4j drivers from 1.7 to 4.0.

The 4.0 drivers have been designed to work with Neo4j 4.0. In 4.0, the drivers are built to provide a user-friendly and unified API across all languages, to take advantage of all new features and services introduced in Neo4j 4.0.

In previous versions of Neo4j, client-server communication used encrypted local connections and generated a self-signed certificate out of the box. In 4.0 however, the default is set to unencrypted. Please see [Driver Manual ▯ Connection URIs](#) for more information.

Neo4j 4.0 introduces a reactive API, compatible with the [Reactive Streams](#) standard. This enables fine-grained control of the data flow for Cypher query results, including the ability to pause or cancel part-way through. Read more in [Driver Manual ▯ Queries and results](#).

When using the 4.0 driver to connect to a 4.0 database, it is possible to work with multiple databases. From a driver API perspective, this means that one database must be selected for use as an execution context for transactions within a session. This can be configured on session construction. If no database is selected, the driver will connect to the server's default database.



Drivers 1.7 work in fallback mode with Neo4j 4.0. They do not support features introduced in Neo4j 4.0, such as [multiple databases](#), [Neo4j Fabric](#), and [fine-grained access control](#). To be able to run multiple databases online concurrently and to do distributed queries over them, you must migrate from 1.7 to 4.0.

The examples in this chapter are mainly written in Java, using the Java driver. However, similar code can be translated to other languages.

6.1. New driver releases

Starting with Neo4j 4.0, the versioning scheme for the database, driver and protocol are all aligned. For supported drivers, this means that the version number will go from 1.7 to 4.0.

The new 4.0 drivers for different languages can be found with the links below:

- [.NET Driver 4.0](#)
- [Java Driver 4.0](#)
- [JavaScript Driver 4.0](#)
- [Python Driver 4.0](#)



The Go Driver 4.0 is under construction.

The current stable (version 1.8) for the Go Driver will work in fallback mode with Neo4j 4.0. Therefore, all functionality that exists in Neo4j 3.5 will also be available in Neo4j 4.0, but new functionality introduced in 4.0 will not.

Note that a 1.7+ driver communicating with a 4.0 server may need to have encryption explicitly switched off. This is due to a change in the defaults between Neo4j 3.x and 4.0.

6.2. Compatibility

The compatibility between Neo4j 3.5 and 4.0, and 4.0 Bolt drivers is illustrated in the tables below:

Table 2. Protocols

	Neo4j 4.0	Neo4j 3.5
Bolt v4.0	All features fully supported.	Not supported.
Bolt v3	All features fully supported, but the support may be removed in next version.	All features fully supported.
Bolt v2	Not supported.	All features fully supported, but the support may be removed in next version.
Bolt v1	Not supported.	All features fully supported, but the support may be removed in next version.

Table 3. Drivers

	Neo4j 4.0		Neo4j 3.5	
	Bolt version	Support	Bolt version	Support
Java Driver 4.0	Bolt v4.0	All features fully supported.	Bolt v3	All features fully supported, but the support may be removed in next version.
.NET Driver 4.0	Bolt v4.0	All features fully supported.	Bolt v3	All features fully supported, but the support may be removed in next version.
JavaScript Driver 4.0	Bolt v4.0	All features fully supported.	Bolt v3	All features fully supported, but the support may be removed in next version.
Python Driver 4.0	Bolt v4.0	All features fully supported.	Bolt v3	All features fully supported, but the support may be removed in next version.
Go Driver 1.8¹	Bolt v3	All features partially supported.	Bolt v3	All features fully supported.
Go Driver 1.7¹	Bolt v3	All features partially supported.	Bolt v3	All features fully supported.

^[1]Neo4j The Go Driver 4.0 is still under construction. Please refer to <https://github.com/neo4j/neo4j-go-driver> for the latest versions that are available.

6.3. What's new?

- Bolt v4.0 is implemented in both 4.0 drivers and 4.0 servers.
- Reactive API is now available with 4.0 servers. To make use of the reactive API, the starting point is `RxSession` on the driver object.
- With 4.0 servers, session instances should now be acquired against a specific database. Causal chaining is still respected on each database (transactions cannot span across multiple databases). The driver itself connects to Neo4j DBMS.
- A new feature detection method `driver.supportsMultiDb()` is added for querying whether the remote database supports multiple databases.
- A new `driver.verifyConnectivity()` method is introduced for connectivity verification purposes. The driver instances by default will not verify DBMS availability after construction.
- New connection URI schemes with variants that contain extra encryption and trust information - `neo4j+s`, `bolt+s`, `neo4j+ssc` and `bolt+ssc`. The `+s` variants enable encryption with a full certificate check, and the `+ssc` variants enable encryption, but with no certificate check. This latter variant is designed specifically for use with self-signed certificates. For more information, see [Additional URI Schemes](#).

Example 4. Detecting multiple database support

```
import org.neo4j.driver.Driver;
import org.neo4j.driver.Result;
import org.neo4j.driver.Session;
import org.neo4j.driver.SessionConfig;
import org.neo4j.driver.Values;

...
private final Driver driver;
...

public void printGreeting( final String message )
{
    SessionConfig sessionConfig = driver.supportsMultiDb() ? SessionConfig.forDatabase( "neo4j" )
                                                            : SessionConfig.defaultConfig();

    try ( Session session = driver.session( sessionConfig ) )
    {
        String greeting = session.writeTransaction( tx -> {
            Result result = tx.run( "CREATE (a:Greeting) SET a.message = $message RETURN a.message +
', from node ' + id(a)",
                                Values.parameters( "message", message ) );
            return result.single().get( 0 ).asString();
        } );
        System.out.println( greeting );
    }
}
```

6.4. Breaking changes

- The driver's default configuration for encrypted is now false (meaning that driver will only attempt plain text connections by default). Connections to encrypted services (such as Neo4j Aura) should now explicitly be set to encrypted.
- When encryption is explicitly enabled, the default trust mode is to trust the CAs that are trusted by operating system. This means that encrypted connections to servers holding self-signed certificates will now fail on certificate verification by default.
- Hostname verification is turned on by default when encryption is turned on.
- `v1` is removed from drivers' package name. For example, in the Java driver, all public APIs are in the package `org.neo4j.driver` instead of the old `org.neo4j.driver.v1`.
- The `neo4j://` scheme replaces `bolt+routing://` and can be used for both clustered and single-instance configurations. This is a rename only, and `neo4j://` URIs can still be used to communicate with Neo4j 3.x clusters. Please note though that Neo4j 3.x standalone instances do not expose a routing interface.

The `bolt://` scheme is used for direct connection to a particular Neo4j server. This scheme is no longer required for standalone machines, however. Neo4j 4.0 now exposes a routing interface for all deployment topologies, allowing `neo4j://` URIs to be used for all deployments. The `bolt://` scheme is now mainly only useful when targeting a specific machine, rather than an entire service. This can be a certain server in a Causal Cluster or the one server in a single-instance environment.

- For drivers where synchronous and asynchronous methods are both implemented, asynchronous methods have been extracted out and put in `AsyncSession`, whereas synchronous methods remain in `Session`. This change ensures that blocking and non-blocking APIs can never be mixed together.
- `Driver#session` method now makes use of a session configuration object or option builder, rather than method arguments.
- Bookmark has changed from a string, and/or a list of strings, to a Bookmark object.
- For synchronous Transaction API, `Transaction#success` and `Transaction#failure` have been removed.

The `success/close` pattern for Transaction objects is now obsolete and has been fully superseded by `commit` and `rollback` methods. However, unlike `Transaction#success`, which only marks the transaction to be successful and then waits for `Transaction#close` to actually perform the real commit, `Transaction#commit` commits the transaction immediately.

A transaction in 4.0 can only be committed or rolled back once. If a transaction is not committed explicitly using `Transaction#commit`, `Transaction#close` will roll back the transaction.

- `Statement` has been renamed to `Query`. `StatementResult` has been renamed to `Result`. Similarly, `StatementResultCursor` has been renamed to `ResultCursor`.
- A result can only be consumed once.

A result is consumed if either the query result has been discarded by invoking `Result#consume`, and/or the outer scope where the result is created, such as a transaction or a session, has been closed. Attempts to access consumed results will be responded with a `ResultConsumedException`.

- The experimental `StatementRunner.typeSystem()` has moved to `Driver.defaultTypeSystem()`.
- `LoadBalancingStrategy` is removed from Config class, and the drivers always default to `LeastConnectedStrategy`.
- The recommended Driver Connection URI scheme is as follows:

Table 4. Recommended Driver Connection URI scheme.

		4.0 drivers	1.7 drivers
4.0 Neo4j	Single instance	neo4j	bolt
	Cluster core members	neo4j	neo4j (bolt+routing)
	Cluster read replicas	neo4j	bolt
3.5 Neo4j	Single instance	bolt	bolt
	Cluster core members	neo4j	neo4j (bolt+routing)
	Cluster read replicas	bolt	bolt

6.4.1. Driver-specific breaking changes

In addition to the breaking changes mentioned above, which apply in general for all drivers, the following drivers have further breaking API changes:

.NET Driver

- The `Neo4j.Driver` package contains only the asynchronous API.
 - Synchronous session API (Simple API) has been moved to the `Neo4j.Driver.Simple` package.
 - Reactive API is presented in the `Neo4j.Driver.Reactive` package.
- The `IDriverLogger` has been renamed to `ILogger`.
- `TrustStrategy` is replaced with `TrustManager`.

See full changelog: <https://github.com/neo4j/neo4j-dotnet-driver/wiki/4.0-changelog>

Example 5. Migrating from the 1.7 .NET Driver to the 4.0 .NET Driver.

Example code for the 4.0 .NET driver

```
using Neo4j.Driver.Simple;
...
private readonly IDriver _driver;
private readonly string
_previousNeo4jSessionBookmark;
...
public void PrintGreeting(string message)
{
    using (ISession session = _driver.Session(o
=>
        o.WithDatabase("neo4j")
        .WithDefaultAccessMode(AccessMode.Write)
        .WithBookmarks
(_previousNeo4jSessionBookmark)))
    {
        using (ITransaction transaction =
session.BeginTransaction())
        {
            Query query = new Query("CREATE
(a:Greeting) SET a.message = $message RETURN
a.message + ', from node ' + id(a)", new
Dictionary<string, object>{{"message",
message}});
            IResult result = transaction.Run
(query);

            string greeting = result.
Single()[0].As<string>();
            Console.WriteLine(greeting);
            transaction.Commit(); // commit
immediately here
        }
        _previousNeo4jSessionBookmark =
session.LastBookmark;
    }
}
```

Example code for the 1.7 .NET driver

```
using Neo4j.Driver;
...
private readonly IDriver _driver;
private readonly string
_previousSessionBookmark;
...
public void PrintGreeting(string message)
{
    using (ISession session = _driver.Session(
        AccessMode.Write,
        _previousSessionBookmark))
    {
        using (ITransaction transaction =
session.BeginTransaction())
        {
            Statement query = new Statement
("CREATE (a:Greeting) SET a.message = $message
RETURN a.message + ', from node ' + id(a)", new
Dictionary<string, object>{{"message",
message}});
            IStatementResult result =
transaction.Run(query);
            transaction.Success(); // mark
success, actually commit will happen in
transaction.Dispose()
            var greeting = result.Single()[0].As
<string>();
            Console.WriteLine(greeting);
        }
        _previousSessionBookmark = session
.LastBookmark;
    }
}
```

JavaScript Driver

- `session#close()` and `driver#close()` both now return `Promises`, and no longer accept callback function arguments.
- `driver.onError` and `driver.onCompleted` callbacks have been completely removed. Errors should be monitored on related code paths (i.e. through `Promise#catch`, etc.).

See full changelog: <https://github.com/neo4j/neo4j-javascript-driver/wiki/4.0-changelog>

Example 6. Migrating from the 1.7 JavaScript Driver to the 4.0 JavaScript Driver.

Example code for the 4.0 JavaScript driver	Example code for the 1.7 JavaScript driver
<pre>var neo4j = require('neo4j-driver') ... const driver = neo4j.driver(uri, neo4j.auth .basic(user, password)) ... const session = driver.session() try { const tx = session.beginTransaction() const result = await tx.run('CREATE (a:Greeting) SET a.message = \$message RETURN a.message + ", from node " + id(a)', { message: 'hello, world' }) const greeting = result.records[0].get(0) console.log(greeting) await tx.commit() } finally { await session.close() }</pre>	<pre>var neo4j = require('neo4j-driver').v1 ... const driver = neo4j.driver(uri, neo4j.auth .basic(user, password)) ... const session = driver.session() try { const tx = session.beginTransaction() const result = await tx.run('CREATE (a:Greeting) SET a.message = \$message RETURN a.message + ", from node " + id(a)', { message: 'hello, world' }) const greeting = result.records[0].get(0) console.log(greeting) await tx.commit() } finally { session.close(callback) // another session can be chained in callback }</pre>

Java Driver

See full changelog: <https://github.com/neo4j/neo4j-java-driver/wiki/4.0-changelog>

Example 7. Migrating from the 1.7 Java Driver to the 4.0 Java Driver.

Example code for the 4.0 Java Driver

```
import org.neo4j.driver.Bookmark;
import org.neo4j.driver.Driver;
import org.neo4j.driver.Query;
import org.neo4j.driver.Result;
import org.neo4j.driver.Session;
import org.neo4j.driver.SessionConfig;
import org.neo4j.driver.Transaction;
import org.neo4j.driver.Values;
...

private final Driver driver;
...
public void printGreeting( String message,
Bookmark bookmark )
{
    SessionConfig sessionConfig = SessionConfig
        .builder()
        .withDatabase( "neo4j" )
        .withDefaultAccessMode( AccessMode.WRITE )
        .withBookmarks( bookmark ).build();

    try ( Session session = driver.session(
        sessionConfig );
        Transaction transaction = session
        .beginTransaction() )
    {
        Query query = new Query( "CREATE
(a:Greeting) SET a.message = $message RETURN
a.message + ', from node ' + id(a)", Values
        .parameters( "message", message ) );

        Result result = transaction.run( query );
        String greeting = result.single().get( 0
        ).asString();
        System.out.println( greeting );
        transaction.commit(); // commit
        immediately here
    }
}
```

Example code for the 1.7 Java Driver

```
import org.neo4j.driver.v1.AccessMode;
import org.neo4j.driver.v1.Driver;
import org.neo4j.driver.v1.Session;
import org.neo4j.driver.v1.StatementResult;
import org.neo4j.driver.v1.Transaction;
import org.neo4j.driver.v1.Values;
...

private final Driver driver;
...
public void printGreeting( String message,
String bookmark )
{

    try ( Session session = driver.session(
        AccessMode.WRITE, bookmark );
        Transaction transaction = session
        .beginTransaction() )
    {
        Statement query = new Statement( "CREATE
(a:Greeting) SET a.message = $message RETURN
a.message + ', from node ' + id(a)", Values
        .parameters( "message", message ) );

        StatementResult result = transaction.
run( query );
        transaction.success(); // mark success,
        actually commit will happen in
        transaction.close()
        String greeting = result.single().get( 0
        ).asString();
        System.out.println( greeting );
    }
}
```

Python Driver

- Renamed Configuration `max_retry_time` is renamed to `max_transaction_retry_time`.
- Renamed Configuration `access_mode` to `default_access_mode` and this is now a **keyword argument**.
- Renamed Exception `neo4j.exceptions.CypherError` to `neo4j.exceptions.Neo4jError`.
- Removed Exception `neo4j.exceptions.ConnectionExpired`.
- Removed `transaction.success` flag.
- `neo4j.Record` helper function `Result.value(key=0, default=None)`, the `key` parameter have changed name from `item` to `key`.
- `neo4j.Record` helper function `Result.values(*keys)`, the `*keys` paramatere have changed name from `*items` to `*keys`.
- `neo4j.Record` helper function `Result.data(*keys)`, the `*keys` paramatere have changed name from `*items` to `*keys`.
- `neobolt` is not a dependency any more.
- `neotime` is not a dependency any more.

- `pytz` is a dependency now.
- `Transaction.sync()` has been removed. Use `Result.consume()` if the behaviour is to exhaust the result object.
- `Transaction.success` has been removed.
- `Transaction.close()` behaviour changed. Will now only perform rollback if no commit have been performed.
- `Session.sync()` has been removed. Use `Result.consume()` if the behaviour is to exhaust the result object.
- `Session.detach()` has been removed. Use `Result.consume()` if the behaviour is to exhaust the result object.
- `Session.next_bookmarks()` has been removed.
- `Session.has_transaction()` has been removed.
- `Session.closed()` has been removed.

See full changelog: <https://github.com/neo4j/neo4j-python-driver/wiki/4.0-changelog>

6.5. Back-pressure

Neo4j 4.0 introduces client-side back-pressure. The concept of client-side back-pressure is as follows; the client will communicate with the remote server regarding how much data it is able to process, and will only request additional data when it is ready to consume more.

The back-pressure concept is naturally compatible with [Reactive programming](#). As a result, Reactive API support is added into all language drivers in 4.0 driver releases.

The Java driver Reactive API exposes a raw Publisher-Subscriber API, which is defined by [reactive streams](#). When using Java driver's Reactive API, it is anticipated that it is used with a reactive library, such as [Project Reactor](#) and/or [RxJava](#).

For the .NET and JavaScript drivers Reactive API, the drivers are already supplied with reactive libraries. The built-in [System.Reactive](#) has been used in .NET driver. As for JavaScript driver, the popular [RxJs](#) library is adopted. These two libraries, as well as Rxjava all belong to the same reactive framework [ReactiveX](#).

To use the drivers' Reactive API, a preliminary knowledge of reactive programming is necessary. Details of how to use Neo4j Reactive Driver API can be found in the [Neo4j Driver Manual](#).

However, back-pressure is not only limited to the driver's Reactive APIs. All other APIs, such as simple and async, by default have back-pressure enabled when handling query execution results.

Table 5. How back-pressure is implemented in the different language driver session APIs

	Simple API	Async API	Reactive API
Java driver	Record buffer	Record buffer	Raw Publisher-Subscriber API
.NET driver	Record buffer	Record buffer	Record buffer
Javascript driver	Not applicable	Record buffer	Record buffer

6.5.1. Back-pressure with Bolt v4.0 and record buffer

The Neo4j 4.0 server and drivers implement Bolt v4.0. One of the main features introduced in this Bolt version is pulling query results (records) in batches. In previous Bolt versions, the complete result set is always pulled in one batch from a server to a driver. Bolt v4.0 enables us to pull these results in

multiple batches where the size of each can be defined by a `fetchSize`. By default, the drivers use a `fetchSize` of 1000 records.

With the introduction of batching of records, drivers could now implement client-side back-pressure. For each result, the driver keeps a record buffer of unconsumed records. The buffer size is the same as `fetchSize` for each batch. The pulling of records from the server will be paused when the buffer is more than 70% full, and the record pulling is re-enabled once the buffer is less than 30% full. With the default `fetchSize` of 1000 records, record pulling is paused when more than 700 records are in the buffer and is resumed when the buffer drops below 300.

Example 8. Set default `fetchSize` on driver and alter the default value on session.

```
import org.neo4j.driver.AuthTokens;
import org.neo4j.driver.Config;
import org.neo4j.driver.Driver;
import org.neo4j.driver.GraphDatabase;
import org.neo4j.driver.Session;
import org.neo4j.driver.SessionConfig;
...

Config config = Config.builder().withFetchSize( 2000 ).build();
Driver driver = GraphDatabase.driver( uri, AuthTokens.basic( user, password ), config );

SessionConfig sessionConfig = SessionConfig.builder()
    .withDatabase( "neo4j" )
    .withFetchSize( 100 )
    .build();

try ( Session session = driver.session( sessionConfig ) ) {...}
```

6.5.2. Java driver Reactive API

The Java driver's Reactive API exposes a very low level Publisher-Subscriber API. As a result, it will not perform any kind of back-pressure by default. Instead, we expect driver users to make use of a reactive framework to utilise back-pressure. Depending on the choice of the reactive framework, the framework may apply back-pressure by pausing the pulling of the data from a Neo4j server, or dropping data when there is too much to process.

6.6. Multiple databases

With the addition of multiple databases in 4.0, you can now specify which database to work with. When constructing a session you can specify in the session configuration which database the session is linked to. Queries will then be executed against that database for the duration of the session. Not specifying a database will result in the session being linked to the default database as specified in the server configuration, see [Operations manual ▯ The default database](#). When using 4.0 drivers with 4.0 Neo4j Servers, we always recommend to specify the database of each session explicitly.

Example 9. Selecting a database for a session.

```
import org.neo4j.driver.Driver;
import org.neo4j.driver.Session;
import org.neo4j.driver.SessionConfig;
...

try ( Session session = driver.session( SessionConfig.forDatabase( "neo4j" ) ) ) {...}
```

While managing multiple databases is primarily a feature for Neo4j Enterprise Edition, users of Neo4j Community Edition will still need to use the `system` database when carrying out administrative operations on the database. See [Operations manual ▯ The system database](#) and [Cypher manual ▯ Administration](#) for more information.

6.6.1. Bookmarks



Bookmarks are generally handled internally by the driver. Applications typically only need to work with bookmarks directly when chaining sessions.

When using bookmarks in a multiple database context, the base rule is that bookmarks can only be passed among sessions for the same database. This is because the bookmarks (and/or transactions) cannot cross multiple databases in Neo4j 4.0. There is one exception however, the bookmarks generated by the `system` database can be used with other databases.

Example 10. Using system bookmark with another database to ensure the updated system status.

```
import org.neo4j.driver.Bookmark;
import org.neo4j.driver.Driver;
import org.neo4j.driver.Result;
import org.neo4j.driver.Session;
import org.neo4j.driver.SessionConfig;
...

Bookmark sysBookmark;
try ( Session session = driver.session( SessionConfig.forDatabase( "system" ) ) )
{
    session.writeTransaction( tx -> {
        Result result = tx.run( "CREATE database foo" );
        return result.consume();
    } );
    sysBookmark = session.lastBookmark();
}

try ( Session session = driver.session( SessionConfig.builder().withDatabase( "foo" ).withBookmarks(
    sysBookmark ).build() ) )
{
    session.writeTransaction( tx -> {
        Result result = tx.run( "CREATE (n)" );
        return result.consume();
    } );
}
```

6.7. Configure SSL Policy for Bolt server and HTTPS server

Neo4j 3.5 always allows encrypted connections with the default configuration. In case no certificate is installed before a server starts, self-signed certificates will be automatically generated. However, in 4.0 the default encryption setting is off and Neo4j will no longer generate certificates when none is provided. As a result, Bolt server only allows plaintext connections, and HTTPS server is not enabled by default. The table below summarizes the default behaviour change between 3.5 and 4.0 regarding encryption and certificates.

Table 6. Encryption and certificates differences between 3.5 and 4.0 servers.

	3.5 Neo4j Bolt Server	4.0 Neo4j Bolt Server	3.5 Neo4j HTTPS Server	4.0 Neo4j HTTPS Server
Server Enabled	Yes	Yes	Yes	No
Encryption on client connections	Optional	Not allowed	Always	Always
Certificates	Auto-generated self-signed certificates if not provided	None	Auto-generated self-signed certificates if not provided	None

	3.5 Neo4j Bolt Server	4.0 Neo4j Bolt Server	3.5 Neo4j HTTPS Server	4.0 Neo4j HTTPS Server
Default Certificates Path	<code>\$neo4jHome/certificates</code>	None	<code>\$neo4jHome/certificates</code>	None
Default Certificate Names	<code>neo4j.key</code> <code>neo4j.cert</code>	<code>private.key</code> <code>public.crt</code>	<code>neo4j.key</code> <code>neo4j.cert</code>	<code>private.key</code> <code>public.crt</code>

In order to re-enable encryption in 4.0, we need to configure the SSL policy in the Neo4j config file. Given certificates named `public.crt` and `private.key` in folder `$neo4jHome/certificates/bolt` for Bolt server, and certificates with the same file names in folder `$neo4jHome/certificates/https` for HTTPS server. The example below shows how to turn encryption back on for the Bolt server and re-enable the HTTPS server.

Example 11. Turn encryption on for Bolt v4.0 server.

```
dbms.connector.bolt.enabled=true
dbms.connector.bolt.tls_level=OPTIONAL # allows both encrypted and unencrypted driver
connections

dbms.ssl.policy.bolt.enabled=true
dbms.ssl.policy.bolt.base_directory=certificates/bolt
#dbms.ssl.policy.bolt.private_key=private.key # Optional if the file name is the same as the
default.
#dbms.ssl.policy.bolt.public_certificate=public.crt # Optional if the file name is the same as
the default.
```

Example 12. Enable the HTTPS 4.0 server.

```
dbms.connector.https.enabled=true

dbms.ssl.policy.https.enabled=true
dbms.ssl.policy.https.base_directory=certificates/https
#dbms.ssl.policy.https.private_key=private.key # Optional if the file name is the same as the
default.
#dbms.ssl.policy.https.public_certificate=public.crt # Optional if the file name is the same as
the default.
```

6.8. Additional URI Schemes

Since **v4.0.1** of the Java and .NET drivers, and **v4.0.2** of the JavaScript driver, you are able to configure the encryption and trust settings of the driver directly through the connection URI.

The `neo4j+s` and `bolt+s` schemes enable encryption and full certificate checks against the system's local CA store. The `neo4j+ssc` and `bolt+ssc` schemes also enable encryption with no certificate checks, typically for use with self-signed certificates.

Table 7. Available URIs

URI	Routing	Description
<code>neo4j</code>	Yes	Unsecured

URI	Routing	Description
neo4j+s	Yes	Secured with full certificate
neo4j+ssc	Yes	Secured with self-signed certificate
bolt	No	Unsecured
bolt+s	No	Secured with full certificate
bolt+ssc	No	Secured with self-signed certificate

Using these new URI schemes is not compatible with configuring encryption and trust with the Configuration API. Otherwise, this does not effect the behaviour of the existing `neo4j` and `bolt` schemes.

For more information, see [Driver Manual ▯ Connection URIs](#).

Chapter 7. Tutorial: Back up and restore a database

This tutorial provides a detailed example of how to back up and restore a 3.5 database on a running 4.x standalone instance.

The following example assumes that your database has users and roles associated with it and describes how to back it up, upgrade it, and then restore it on a running standalone instance.

7.1. Prepare to back up the database

Before you perform the backup, it is good to take a note of the data and metadata of the database that you want to restore. You can use this information later to verify that the restore is successful and to recreate the database users and roles. In this example, the database uses the Movie Graph dataset from the Neo4j Browser ▢ Favorites ▢ Example Graphs.

1. In the 3.5 Neo4j instance, where the database is running, navigate to the `/bin` folder and log in to the Cypher Shell command-line console with your credentials. For more information about the Cypher Shell command-line interface (CLI) and how to use it, see [Operations Manual ▢ Cypher Shell](#).

```
./cypher-shell -u neo4j -p <password>
```

```
Connected to Neo4j at neo4j://localhost:7687 as user neo4j.  
Type :help for a list of available commands or :exit to exit the shell.  
Note that Cypher queries must end with a semicolon.
```

2. Run a query to count the number of nodes in the database.

```
MATCH (n) RETURN count(n) AS countNode;
```

```
+-----+  
| countNode |  
+-----+  
| 171      |  
+-----+
```

```
1 row available after 22 ms, consumed after another 1 ms
```

3. Run a query to count the number of relationships.

```
MATCH (n)-[r]->(n) RETURN count(r) AS countRelationships;
```

```
+-----+  
| countRelationships |  
+-----+  
| 253                |  
+-----+
```

```
1 row available after 29 ms, consumed after another 0 ms
```

4. Run the following two queries to see if there is a schema defined.

```
CALL db.constraints()
```

```
0 rows available after 2 ms, consumed after another 0 ms
```

The result shows that there are no constraints defined.

```
CALL db.indexes;
```

```
+-----+
+-----+
| description          | indexName | tokenNames | properties | state   | type          |
| progress | provider          | id | failureMessage |
+-----+
+-----+
| "INDEX ON :Movie(title)" | "index_1" | ["Movie"] | ["title"] | "ONLINE" | "node_label_property" |
| 100.0 | {version: "1.0", key: "native-btree"} | 1 | "" |
+-----+
+-----+
```

The result shows that there is one index defined on the property **title** of the **:Movie** node.

5. Run a query to list all users associated with this database and their roles.

```
CALL dbms.security.listUsers;
```

```
+-----+
| username | roles          | flags |
+-----+
| "user1"  | ["editor", "reader"] | []    |
| "neo4j"  | ["admin"]       | []    |
+-----+
```

```
2 rows available after 2 ms, consumed after another 0 ms
```

The result shows two users - the default **neo4j** user, which has **admin** privileges, and a custom user **user1**, which has the combined privileges of the built-in roles **editor** and **reader**.

6. Exit the Cypher Shell command-line console.

```
:exit;
Bye!
```

7.2. Back up the database

Now you are ready to back up the database.

Navigate to the **/bin** folder, and run the following command to back up the database in your targeted folder. If the folder where you want to place your backup does not exist, you have to create it. In this example, it is called **/tmp/3.5.24**.

```
./neo4j-admin backup --backup-dir=/tmp/3.5.24 --name=graphdbbackup
```

For details on performing a backup and the different command options, see [Operations Manual](#) □

Perform a backup.

7.3. Restore the database backup on a 4.x standalone instance

You have a running Neo4j 4.x instance, and you want to restore your backed up database in it.

1. In the *neo4j.conf* file of the 4.x standalone instance, set `dbms.allow_upgrade=true`.



If your Neo4j standalone instance is a version earlier than 4.1, you have to restart it for the configuration to take effect.

2. Navigate to the */bin* folder and run the following command to restore the database backup.

```
./neo4j-admin restore --from=/tmp/3.5.24/graphdbbackup --database=graphdbbackup
```

3. Run the following command to verify that the database `graphdbbackup` exists:

```
ls -al ../data/databases
```

```
total 0
drwxr-xr-x@ 6 username staff 192 4 Dec 14:15 .
drwxr-xr-x@ 5 username staff 160 7 Dec 09:35 ..
drwxr-xr-x 42 username staff 1344 4 Dec 14:15 graphdbbackup
drwxr-xr-x 37 username staff 1184 4 Dec 14:06 neo4j
-rw-r--r-- 1 username staff 0 4 Dec 14:06 store_lock
drwxr-xr-x 38 username staff 1216 4 Dec 14:06 system
```

However, restoring a database does not automatically create it. Therefore, it will not be visible if you do `SHOW DATABASES` in Cypher Shell or Neo4j Browser.

4. Log in to the Cypher Shell command-line console.
5. Change the active database to `system (:USE system;)`, and create the `graphdbbackup` database.

```
CREATE DATABASE graphdbbackup;
```

```
0 rows available after 145 ms, consumed after another 0 ms
```

6. Verify that the `graphdbbackup` database is online.

```
SHOW DATABASES;
```



```

+-----+
--+
| name          | address          | role          | requestedStatus | currentStatus | error |
default |
+-----+
--+
| "graphdbbackup" | "localhost:7687" | "standalone" | "online"        | "online"      | ""    | FALSE
| "neo4j"         | "localhost:7687" | "standalone" | "online"        | "online"      | ""    | TRUE
| "system"        | "localhost:7687" | "standalone" | "online"        | "online"      | ""    | FALSE
+-----+
--+

3 rows available after 175 ms, consumed after another 9 ms

```

- Change the active database to **graphdbbackup**, and repeat steps 2, 3, and 4 from the section [Prepare to back up the database](#), to verify that all the data has been successfully restored.

7.4. Recreate the database users and roles

You manually recreate all users and roles of the restored database using your notes from step 5 of the section [Prepare to back up the database](#) and the [Cypher Manual](#) [Cypher administration commands](#).

*Example 13. Run the following commands against the **system** database to recreate the **graphdbbackup** database's custom users and roles.*

- Create the user **user1**.

```

CREATE USER user1 IF NOT EXISTS
SET PASSWORD 'password'
SET STATUS ACTIVE;

```

- Grant the role **reader** to the user **user1**.

```

GRANT ROLE reader TO user1;

```

- Grant the role **editor** to the user **user1**.

```

GRANT ROLE editor TO user1;

```

- Verify that the user **user1** has the same roles as in the database backup.

```

SHOW USERS;

```

```

+-----+
| user  | roles                  | passwordChangeRequired | suspended |
+-----+
| "neo4j" | ["admin"]             | FALSE                  | FALSE    |
| "user1" | ["editor", "reader"] | TRUE                   | FALSE    |
+-----+

```

Chapter 8. Tutorial: Back up and restore a database in a Causal Cluster

This tutorial provides a detailed example of how to back up and restore a 3.5 database in a running 4.x Causal Cluster.

The following example assumes that your database has users and roles associated with it and describes how to back it up, upgrade it, and then restore it in a running Causal Cluster. For more information on how to set up a Causal Cluster, see [Operations Manual - Set up a local Causal Cluster](#).



In a Neo4j DBMS, every database is backed up individually. Therefore, it is very important to plan your backup strategy for each of them. For more detailed information on how to design an appropriate backup strategy for your setup, see [Operations Manual - Backup](#).

8.1. Prepare to back up the database

Before you perform the backup, it is good to take a note of the data and metadata of the database that you want to restore. You can use this information later to verify that the restore is successful and to recreate the database users and roles. In this example, the database uses the Movie Graph dataset from the Neo4j Browser ▢ Favorites ▢ Example Graphs.

1. In the 3.5 Neo4j instance, where the database is running, navigate to the `/bin` folder and log in to the Cypher Shell command-line console with your credentials. For more information about the Cypher Shell command-line interface (CLI) and how to use it, see [Operations Manual ▢ Cypher Shell](#).

```
./cypher-shell -u neo4j -p <password>
```

```
Connected to Neo4j at neo4j://localhost:7687 as user neo4j.  
Type :help for a list of available commands or :exit to exit the shell.  
Note that Cypher queries must end with a semicolon.
```

2. Run a query to count the number of nodes in the database.

```
MATCH (n) RETURN count(n) AS countNode;
```

```
+-----+  
| countNode |  
+-----+  
| 171       |  
+-----+
```

```
1 row available after 22 ms, consumed after another 1 ms
```

3. Run a query to count the number of relationships.

```
MATCH (n)-[r]->() RETURN count(r) AS countRelationships;
```

```
+-----+
| countRelationships |
+-----+
| 253                |
+-----+
```

1 row available after 29 ms, consumed after another 0 ms

- Run the following two queries to see if there is a schema defined.

```
CALL db.constraints()
```

0 rows available after 2 ms, consumed after another 0 ms

The result shows that there are no constraints defined.

```
CALL db.indexes;
```

```
+-----+
+-----+
| description          | indexName | tokenNames | properties | state   | type           |
| progress | provider          |           | id | failureMessage |           |
+-----+
+-----+
| "INDEX ON :Movie(title)" | "index_1" | ["Movie"] | ["title"] | "ONLINE" | "node_label_property" |
| 100.0    | {version: "1.0", key: "native-btree"} | 1 | ""           |           |
+-----+
+-----+
```

The result shows that there is one index defined on the property **title** of the **:Movie** node.

- Run a query to list all users associated with this database and their roles.

```
CALL dbms.security.listUsers;
```

```
+-----+
| username | roles          | flags |
+-----+
| "user1"  | ["editor", "reader"] | []    |
| "neo4j"  | ["admin"]       | []    |
+-----+
```

2 rows available after 2 ms, consumed after another 0 ms

The result shows two users - the default **neo4j** user, which has **admin** privileges, and a custom user **user1**, which has the combined privileges of the built-in roles **editor** and **reader**.

- Exit the Cypher Shell command-line console.

```
:exit;
Bye!
```

8.2. Back up the database

Now you are ready to back up the database.

Navigate to the `/bin` folder, and run the following command to back up the database in your targeted folder. If the folder where you want to place your backup does not exist, you have to create it. In this example, it is called `/tmp/3.5.24`.

```
./neo4j-admin backup --backup-dir=/tmp/3.5.24 --name=graphdbbackup
```

For details on performing a backup and the different command options, see [Operations Manual](#) [Perform a backup](#).

8.3. Migrate the database backup to a 4.x standalone instance

To migrate the database backup and upgrade its store, you can spin up a standalone Neo4j instance and use the command `neo4j-admin restore`.

1. In the `neo4j.conf` file of a 4.x standalone instance, set `dbms.allow_upgrade=true`.
2. Start the server.
3. Navigate to the `/bin` folder and run the following command to restore the database backup.

```
./neo4j-admin restore --from=/tmp/3.5.24/graphdbbackup --database=graphdbbackup
```

4. Run the following command to verify that the database `graphdbbackup` exists:

```
ls -al ../data/databases
```

```
total 0
drwxr-xr-x@ 6 username staff 192 4 Dec 14:15 .
drwxr-xr-x@ 5 username staff 160 7 Dec 09:35 ..
drwxr-xr-x 42 username staff 1344 4 Dec 14:15 graphdbbackup
drwxr-xr-x 37 username staff 1184 4 Dec 14:06 neo4j
-rw-r--r-- 1 username staff 0 4 Dec 14:06 store_lock
drwxr-xr-x 38 username staff 1216 4 Dec 14:06 system
```

However, restoring a database does not automatically create it. Therefore, it will not be visible if you do `SHOW DATABASES` in Cypher Shell or Neo4j Browser.

5. Log in to the Cypher Shell command-line console.
6. Change the active database to `system (:USE system;)`, and create the `graphdbbackup` database.

```
CREATE DATABASE graphdbbackup;
```

```
0 rows available after 145 ms, consumed after another 0 ms
```

7. Verify that the `graphdbbackup` database is online.

```
SHOW DATABASES;
```

```

+-----+
--+
| name          | address          | role          | requestedStatus | currentStatus | error |
default |
+-----+
--+
| "graphdbbackup" | "localhost:7687" | "standalone" | "online"        | "online"      | ""    | FALSE
|
| "neo4j"         | "localhost:7687" | "standalone" | "online"        | "online"      | ""    | TRUE
|
| "system"        | "localhost:7687" | "standalone" | "online"        | "online"      | ""    | FALSE
|
+-----+
--+

3 rows available after 175 ms, consumed after another 9 ms

```

- Change the active database to **graphdbbackup**, and repeat steps 2, 3, and 4 from the section [Prepare to back up the database](#), to verify that all the data has been successfully restored.

8.4. Back up the migrated database

To restore the migrated database in your cluster, you create a backup of it.

Navigate to the `/bin` folder, and run the following command to back up the database in your targeted folder. If the folder where you want to place your backup does not exist, you have to create it. In this example, it is called `/tmp/4.0.11`.

```
./neo4j-admin backup --backup-dir=/tmp/4.0.11 --database=graphdbbackup
```

For details on performing a backup and the different command options, see [Operations Manual - Back up an online database](#).

Now you are ready to restore your migrated database backup into your running cluster.

8.5. Restore the migrated database on all cluster members

On each cluster member, navigate to the `/bin` folder, and run the following command to restore the backup of the migrated database. For details on performing a restore and the different command options, see [Operations Manual - Restore command](#).

```
./neo4j-admin restore --from=/tmp/4.0.11/graphdbbackup --database=graphdbbackup
```

Then, on each cluster member, run the following command to verify that the database **graphdbbackup** exists:

```
ls -al ../data/databases
```

```

total 0
drwxr-xr-x@ 6 username staff 192 7 Dec 09:41 .
drwxr-xr-x@ 5 username staff 160 4 Dec 14:17 ..
drwxr-xr-x 38 username staff 1216 7 Dec 09:41 graphdbbackup
drwxr-xr-x 37 username staff 1184 7 Dec 08:57 neo4j
-rw-r--r-- 1 username staff 0 4 Dec 14:17 store_lock
drwxr-xr-x 38 username staff 1216 7 Dec 08:56 system

```

8.6. Create the database backup on the cluster leader

You create the database backup **only on one of your cluster members** using the command `CREATE DATABASE`. The command is automatically routed to the leader, and from there, to the other cluster members.

1. In the Cypher Shell command-line console on one of the cluster members, use the `system` database and create the database `graphdbbackup`.

```
CREATE DATABASE graphdbbackup;
```

```
0 rows available after 132 ms, consumed after another 0 ms
```

2. Verify that the `graphdbbackup` database is online on all members.

```
SHOW DATABASES;
```

```
+-----+
+
| name          | address          | role      | requestedStatus | currentStatus | error | default |
+-----+
+
| "graphdbbackup" | "localhost:7689" | "follower" | "online"        | "online"      | ""    | FALSE   |
| "graphdbbackup" | "localhost:7688" | "leader"   | "online"        | "online"      | ""    | FALSE   |
| "graphdbbackup" | "localhost:7687" | "follower" | "online"        | "online"      | ""    | FALSE   |
| "neo4j"         | "localhost:7689" | "leader"   | "online"        | "online"      | ""    | TRUE    |
| "neo4j"         | "localhost:7688" | "follower" | "online"        | "online"      | ""    | TRUE    |
| "neo4j"         | "localhost:7687" | "follower" | "online"        | "online"      | ""    | TRUE    |
| "system"        | "localhost:7689" | "leader"   | "online"        | "online"      | ""    | FALSE   |
| "system"        | "localhost:7688" | "follower" | "online"        | "online"      | ""    | FALSE   |
| "system"        | "localhost:7687" | "follower" | "online"        | "online"      | ""    | FALSE   |
+-----+
+
```

3. Change your active database to `graphdbbackup` and verify that all the data has been successfully restored by completing the steps 2, 3, and 4 of the section [Prepare to back up the database](#).

8.7. Recreate the database users and roles

On all cluster members, manually recreate all users and roles of the restored database using your notes from step 5 of the section [Prepare to back up the database](#) and the [Cypher Manual](#) [Cypher administration commands](#).

Example 14. Run the following commands against the **system** database to recreate the **graphdbbackup** database's custom users and roles.

1. Create the user **user1**.

```
CREATE USER user1 IF NOT EXISTS
SET PASSWORD 'password'
SET STATUS ACTIVE;
```

2. Grant the role **reader** to the user **user1**.

```
GRANT ROLE reader TO user1;
```

3. Grant the role **editor** to the user **user1**.

```
GRANT ROLE editor TO user1;
```

4. Verify that the user **user1** has the same roles as in the database backup.

```
SHOW USERS;
```

```
+-----+-----+-----+-----+
| user   | roles                | passwordChangeRequired | suspended |
+-----+-----+-----+-----+
| "neo4j" | ["admin"]            | FALSE                  | FALSE    |
| "user1" | ["editor", "reader"] | TRUE                   | FALSE    |
+-----+-----+-----+-----+
```

Chapter 9. Tutorial: Back up and copy a database

This tutorial provides a detailed example of how to back up a 3.5 database and use the `neo4j-admin copy` command to migrate it to a 4.x Neo4j standalone instance.

The `neo4j-admin copy` command can be used to clean up database inconsistencies, compact stores, and do a migration at the same time. It allows a large range of migrations, e.g., 3.5.x directly to 4.2.y, skipping the intermediate steps of 3.5.latest → 4.0.latest → 4.1.latest → 4.2.latest. Those steps are needed to migrate the *schema*, but since the `neo4j-admin copy` command does not copy the schema store at all, they are not needed. However, if there is a schema defined, you have to recreate it by running the commands that the `neo4j-admin copy` operation outputs. The `neo4j-admin copy` command can be applied only to an **offline** database.

The following is an example of how to check your 3.5 database store usage, perform a backup, compact the database backup (using `neo4j-admin copy`), and create it in a Neo4j 4.x standalone instance.



Your copied node IDs will be the same, but the relationships will get new IDs. Therefore, if you want to preserve the relationship IDs, follow the regular [backup and restore upgrade path](#).

9.1. Check your 3.5 database store usage

Before you back up and copy your 3.5 database, let's look at the database store usage and see how it changes when you load, delete, and then reload data.

1. Log in to Neo4j Browser of your running 3.5 Neo4j standalone instance, add 100k nodes to the `graph.db` database using the following command:

```
FOREACH (x IN RANGE (1,100000) | CREATE (n:Person {name:x}))
```

2. Create an index on the `name` property of the `Person` node:

```
CREATE INDEX ON :Person(name)
```

3. Use the `dbms.checkpoint()` procedure to flush all cached updates from the page cache to the store files.

```
CALL dbms.checkpoint()
```

4. In your terminal, navigate to the `graph.db` database (`$neo4j_home/data/databases/graph.db`) and run the following command to check the store size of the loaded nodes and properties.

```
ls -alh
```

```
...
-rw-r--r--  1 username  staff   1.4M 26 Nov 15:51 neostore.nodestore.db
-rw-r--r--  1 username  staff   3.9M 26 Nov 15:51 neostore.propertystore.db
...
```


The output reports that the node store (*neostore.nodestore.db*) and the property store (*neostore.propertystore.db*) occupy **1.4M** and **3.9M**, respectively.

5. In Neo4j Browser, delete the nodes created above and run **CALL dbms.checkpoint** again to force a checkpoint.

```
MATCH (n) DETACH DELETE n
```

```
CALL dbms.checkpoint()
```

6. Now, add just one node, force a checkpoint, and repeat step 4 to see if the store size has changed.

```
CREATE (n:Person {name:"John"})
```

```
CALL dbms.checkpoint()
```

If you check the size of the node store and the property store now, they will still be **1.4M** and **3.9M**, even though the database only contains one node and one property. Neo4j does not shrink the store files on the hard drive.



In a production database, where numerous load/delete operations are performed, the result is a significant unused space occupied by store files.

9.2. Back up your 3.5 database

Navigate to the */bin* folder, and run the following command to back up your database in the targeted folder. If the folder where you want to place your backup does not exist, you have to create it. In this example, it is called */tmp/3.5.24*.

```
./neo4j-admin backup --backup-dir=/tmp/3.5.24 --name=graphdbbackup
```

For details on performing a backup and the different command options, see [Operations Manual](#) [Perform a backup](#).

9.3. Copy your 3.5 database backup to 4.x Neo4j

You can use the **neo4j-admin copy** command to reclaim the unused space and create a defragmented copy of your database backup in your 4.x standalone instance.



To speed up the copy operation, you can use the **--from-pagecache** and **--to-pagecache** options to specify how much cache to be allocated when reading the source and writing the destination. As a rule of thumb, **--to-pagecache** should be around 1-2GB, since it mostly does sequential writes. The **--from-pagecache** should then be assigned whatever memory you can spare, since Neo4j does random reads from the source.

1. In your 4.x Neo4j standalone instance, navigate to the */bin* folder and run the following command to create a compacted store copy of your 3.5 database backup. Any inconsistent nodes, properties, and relationships will not be copied over to the newly created store.

```
./neo4j-admin copy --from-path=/private/tmp/3.5.24/graphdbbackup --to-database=compactdb
```

```
Starting to copy store, output will be saved to: $neo4j_home/logs/neo4j-admin-copy-2020-11-26.16.07.19.log
2020-11-26 16:07:19.939+0000 INFO [StoreCopy] ### Copy Data ###
2020-11-26 16:07:19.940+0000 INFO [StoreCopy] Source: /private/tmp/3.5.24/graphdbbackup (page cache 8m)
2020-11-26 16:07:19.940+0000 INFO [StoreCopy] Target: $neo4j_home/data/databases/compactdb (page cache 8m)
2020-11-26 16:07:19.940+0000 INFO [StoreCopy] Empty database created, will start importing readable data from the source.
2020-11-26 16:07:21.661+0000 INFO [o.n.i.b.ImportLogic] Import starting
```

```
Import starting 2020-11-26 16:07:21.699+0000
Estimated number of nodes: 50.00 k
Estimated number of node properties: 50.00 k
Estimated number of relationships: 0.00
Estimated number of relationship properties: 50.00 k
Estimated disk space usage: 2.680MiB
Estimated required memory usage: 8.598MiB
```

```
(1/4) Node import 2020-11-26 16:07:22.220+0000
```

```
Estimated number of nodes: 50.00 k
Estimated disk space usage: 1.698MiB
Estimated required memory usage: 8.598MiB
```

```
..... 5% Δ239ms
..... 10% Δ1ms
..... 15% Δ1ms
..... 20% Δ0ms
..... 25% Δ1ms
..... 30% Δ0ms
..... 35% Δ0ms
..... 40% Δ1ms
..... 45% Δ0ms
..... 50% Δ1ms
..... 55% Δ0ms
..... 60% Δ51ms
..... 65% Δ0ms
..... 70% Δ0ms
..... 75% Δ1ms
..... 80% Δ0ms
..... 85% Δ0ms
..... 90% Δ1ms
..... 95% Δ0ms
..... 100% Δ0ms
```

```
(2/4) Relationship import 2020-11-26 16:07:22.543+0000
```

```
Estimated number of relationships: 0.00
Estimated disk space usage: 1006KiB
Estimated required memory usage: 15.60MiB
```

```
(3/4) Relationship linking 2020-11-26 16:07:22.879+0000
```

```
Estimated required memory usage: 7.969MiB
```

```
(4/4) Post processing 2020-11-26 16:07:23.272+0000
```

```
Estimated required memory usage: 7.969MiB
```

```
..... 5% Δ356ms
..... 10% Δ0ms
..... 15% Δ1ms
..... 20% Δ0ms
..... 25% Δ0ms
..... 30% Δ1ms
..... 35% Δ0ms
..... 40% Δ0ms
..... 45% Δ1ms
..... 50% Δ0ms
..... 55% Δ0ms
..... 60% Δ0ms
..... 65% Δ1ms
..... 70% Δ0ms
..... 75% Δ0ms
..... 80% Δ0ms
..... 85% Δ0ms
..... 90% Δ0ms
..... 95% Δ1ms
..... 100% Δ0ms
```

```
IMPORT DONE in 2s 473ms.
```

```

Imported:
  1 nodes
  0 relationships
  1 properties
Peak memory usage: 15.60MiB
2020-11-26 16:07:24.140+0000 INFO [o.n.i.b.ImportLogic] Import completed successfully, took 2s 473ms.
Imported:
  1 nodes
  0 relationships
  1 properties
2020-11-26 16:07:24.668+0000 INFO [StoreCopy] Import summary: Copying of 100704 records took 4 seconds
(25176 rec/s). Unused Records 100703 (99%) Removed Records 0 (0%)
2020-11-26 16:07:24.669+0000 INFO [StoreCopy] ### Extracting schema ###
2020-11-26 16:07:24.669+0000 INFO [StoreCopy] Trying to extract schema...
2020-11-26 16:07:24.920+0000 INFO [StoreCopy] ... found 1 schema definitions. The following can be
used to recreate the schema:
2020-11-26 16:07:24.922+0000 INFO [StoreCopy]

CALL db.createIndex('index_5c0607ad', ['Person'], ['name'], 'native-btree-1.0', {'spatial.cartesian-
3d.min': [-1000000.0, -1000000.0, -1000000.0], 'spatial.cartesian.min': [-1000000.0,
-1000000.0], 'spatial.wgs-84.min': [-180.0, -90.0], 'spatial.cartesian-3d.max': [1000000.0, 1000000.0,
1000000.0], 'spatial.cartesian.max': [1000000.0, 1000000.0], 'spatial.wgs-84-3d.min': [-180.0, -90.0,
-1000000.0], 'spatial.wgs-84-3d.max': [180.0, 90.0, 1000000.0], 'spatial.wgs-84.max': [180.0, 90.0]})
2020-11-26 16:07:24.923+0000 INFO [StoreCopy] You have to manually apply the above commands to the
database when it is stored to recreate the indexes and constraints. The commands are saved to
$neo4j_home/logs/neo4j-admin-copy-2020-11-26.16.07.19.log as well for reference.

```

2. Run the following command to verify that database has been successfully copied.

```
ls -al ../data/databases
```

```

total 0
drwxr-xr-x@  5 username  staff   160 26 Nov 18:00 .
drwxr-xr-x@  5 username  staff   160 26 Nov 18:00 ..
drwxr-xr-x  35 username  staff  1120 26 Nov 17:58 compactdb
-rw-r--r--   1 username  staff    26 26 Nov 18:00 store_lock
drwxr-xr-x  33 username  staff  1056 26 Nov 18:00 system

```



Copying a database does not automatically create it. Therefore, it will not be visible if you do **SHOW DATABASES** in Cypher Shell or Neo4j Browser.

9.4. Create your compacted backup

You can now create the copied database and compare its store size with the size of the backed up database.

1. Log in to the Cypher Shell command-line console, change the active database to **system** (:USE **system**;) , and create the **compactdb** database. For more information about the Cypher Shell command-line interface (CLI) and how to use it, see [Operations Manual](#) ▯ [Cypher Shell](#).

```
CREATE DATABASE compactdb;
```

```
0 rows available after 145 ms, consumed after another 0 ms
```

2. Verify that the **compactdb** database is online.

```
SHOW DATABASES;
```

```
+-----+
--+
| name          | address          | role          | requestedStatus | currentStatus | error |
default |
+-----+
--+
| "compactdb"    | "localhost:7687" | "standalone"  | "online"        | "online"      | ""    | FALSE
| "neo4j"        | "localhost:7687" | "standalone"  | "online"        | "online"      | ""    | TRUE
| "system"       | "localhost:7687" | "standalone"  | "online"        | "online"      | ""    | FALSE
+-----+
--+

3 rows available after 10 ms, consumed after another 3 ms
```

3. Change your active database to **compactdb** and recreate the schema using the output from the **neo4j-admin copy** command.

```
CALL db.createIndex('index_5c0607ad', ['Person'], ['name'], 'native-btree-1.0', {'spatial.cartesian-3d.min': [-1000000.0, -1000000.0, -1000000.0], 'spatial.cartesian.min': [-1000000.0, -1000000.0], 'spatial.wgs-84.min': [-180.0, -90.0], 'spatial.cartesian-3d.max': [1000000.0, 1000000.0, 1000000.0], 'spatial.cartesian.max': [1000000.0, 1000000.0], 'spatial.wgs-84-3d.min': [-180.0, -90.0, -1000000.0], 'spatial.wgs-84-3d.max': [180.0, 90.0, 1000000.0], 'spatial.wgs-84.max': [180.0, 90.0]});
```

```
+-----+
| name          | labels          | properties    | providerName    | status          |
+-----+
| "index_5c0607ad" | ["Person"]    | ["name"]      | "native-btree-1.0" | "index created" |
+-----+

1 row available after 50 ms, consumed after another 5 ms
```

4. Verify that all the data has been successfully copied. In this example, there should be one node.

```
MATCH (n) RETURN n.name;
```

```
+-----+
| n.name |
+-----+
| "John" |
+-----+

1 row available after 106 ms, consumed after another 2 ms
```

5. Exit the Cypher Shell command-line console.

```
:exit;
Bye!
```

6. Navigate to the **compactdb** database (**\$neo4j_home/data/databases/compactdb**) and check the store size of the copied nodes and properties.

```
ls -alh
```

```
...  
-rw-r--r--  1 username  staff   8.0K 26 Nov 17:58 neostore.nodestore.db  
-rw-r--r--  1 username  staff   8.0K 26 Nov 17:58 neostore.propertystore.db  
...
```

The output reports that the node store and the property store now occupy only **8K** each, compared to the previous **1.4M** and **3.9M**.

Appendix A: Classes removed from public API

This appendix lists the classes that have been removed or excluded from the public API between Neo4j 3.5 and 4.0.

The following table lists classes that have been removed or excluded from the public API:

Classes removed or excluded from the public API
<code>org.neo4j.backup.BackupExtensionService</code>
<code>org.neo4j.backup.BackupTool</code>
<code>org.neo4j.backup.IncrementalBackupNotPossibleException</code>
<code>org.neo4j.backup.OnlineBackupExtensionFactory.Dependencies</code>
<code>org.neo4j.backup.OnlineBackupExtensionFactory</code>
<code>org.neo4j.backup.OnlineBackupKernelExtension.BackupProvider</code>
<code>org.neo4j.backup.OnlineBackupKernelExtension</code>
<code>org.neo4j.backup.OnlineBackupSettings</code>
<code>org.neo4j.backup.TheBackupInterface</code>
<code>org.neo4j.cypher.export.CypherResultSubGraph</code>
<code>org.neo4j.cypher.export.DatabaseSubGraph</code>
<code>org.neo4j.cypher.export.SubGraphExporter</code>
<code>org.neo4j.cypher.export.SubGraph</code>
<code>org.neo4j.graphalgo.CommonEvaluators</code>
<code>org.neo4j.graphalgo.MaxCostEvaluator</code>
<code>org.neo4j.graphdb.DatabaseShutdownException</code>
<code>org.neo4j.graphdb.DependencyResolver.Adapter</code>
<code>org.neo4j.graphdb.DependencyResolver.SelectionStrategy</code>
<code>org.neo4j.graphdb.DependencyResolver</code>
<code>org.neo4j.graphdb.DynamicLabel</code>
<code>org.neo4j.graphdb.DynamicRelationshipType</code>
<code>org.neo4j.graphdb.InvalidTransactionTypeException</code>
<code>org.neo4j.graphdb.PathExpanderBuilder</code>
<code>org.neo4j.graphdb.PathExpanders</code>
<code>org.neo4j.graphdb.PropertyContainer</code>
<code>org.neo4j.graphdb.ResourceUtils</code>
<code>org.neo4j.graphdb.TransactionGuardException</code>
<code>org.neo4j.graphdb.TransientDatabaseFailureException</code>
<code>org.neo4j.graphdb.TransientFailureException</code>
<code>org.neo4j.graphdb.TransientTransactionFailureException</code>
<code>org.neo4j.graphdb.config.BaseSetting</code>
<code>org.neo4j.graphdb.config.InvalidSettingException</code>
<code>org.neo4j.graphdb.config.ScopeAwareSetting</code>

Classes removed or excluded from the public API

org.neo4j.graphdb.config.SettingGroup

org.neo4j.graphdb.config.SettingValidator

org.neo4j.graphdb.event.ErrorState

org.neo4j.graphdb.event.KernelEventHandler.ExecutionOrder

org.neo4j.graphdb.event.KernelEventHandler

org.neo4j.graphdb.event.TransactionEventHandler.Adapter

org.neo4j.graphdb.facade.GraphDatabaseDependencies

org.neo4j.graphdb.facade.GraphDatabaseFacadeFactory.Dependencies

org.neo4j.graphdb.facade.GraphDatabaseFacadeFactory

org.neo4j.graphdb.facade.embedded.EmbeddedGraphDatabase

org.neo4j.graphdb.facade.spi.ClassicCoreSPI

org.neo4j.graphdb.facade.spi.ProcedureGDBFacadeSPI

org.neo4j.graphdb.factory.Description

org.neo4j.graphdb.factory.EditionLocksFactories

org.neo4j.graphdb.factory.EnterpriseGraphDatabaseFactory

org.neo4j.graphdb.factory.GraphDatabaseBuilder.DatabaseCreator

org.neo4j.graphdb.factory.GraphDatabaseBuilder.Delegator

org.neo4j.graphdb.factory.GraphDatabaseBuilder

org.neo4j.graphdb.factory.GraphDatabaseFactoryState

org.neo4j.graphdb.factory.GraphDatabaseFactory

org.neo4j.graphdb.factory.GraphDatabaseSettings.Connector.ConnectorType

org.neo4j.graphdb.factory.GraphDatabaseSettings.Connector

org.neo4j.graphdb.factory.GraphDatabaseSettings.LabelIndex

org.neo4j.graphdb.factory.HighlyAvailableGraphDatabaseFactory

org.neo4j.graphdb.factory.module.DataSourceModule

org.neo4j.graphdb.factory.module.ModularDatabaseCreationContext

org.neo4j.graphdb.factory.module.PlatformModule

org.neo4j.graphdb.factory.module.ProcedureGDSFactory

org.neo4j.graphdb.factory.module.edition.AbstractEditionModule

org.neo4j.graphdb.factory.module.edition.CommunityEditionModule

org.neo4j.graphdb.factory.module.edition.DefaultEditionModule

org.neo4j.graphdb.factory.module.edition.context.DatabaseEditionContext

org.neo4j.graphdb.factory.module.edition.context.DefaultEditionModuleDatabaseContext

org.neo4j.graphdb.factory.module.id.DatabaseIdContext

org.neo4j.graphdb.factory.module.id.IdContextFactoryBuilder

org.neo4j.graphdb.factory.module.id.IdContextFactory

org.neo4j.graphdb.index.AutoIndexer

org.neo4j.graphdb.index.IndexHits

org.neo4j.graphdb.index.IndexManager

org.neo4j.graphdb.index.IndexPopulationProgress

Classes removed or excluded from the public API

org.neo4j.graphdb.index.Index

org.neo4j.graphdb.index.ReadableIndex

org.neo4j.graphdb.index.ReadableRelationshipIndex

org.neo4j.graphdb.index.RelationshipAutoIndexer

org.neo4j.graphdb.index.RelationshipIndex

org.neo4j.graphdb.index.UniqueFactory.UniqueEntity

org.neo4j.graphdb.index.UniqueFactory.UniqueNodeFactory

org.neo4j.graphdb.index.UniqueFactory.UniqueRelationshipFactory

org.neo4j.graphdb.index.UniqueFactory

org.neo4j.graphdb.security.AuthProviderFailedException

org.neo4j.graphdb.security.AuthProviderTimeoutException

org.neo4j.graphdb.security.AuthorizationExpiredException

org.neo4j.graphdb.security.AuthorizationViolationException

org.neo4j.graphdb.security.URLAccessRule

org.neo4j.graphdb.security.URLAccessValidationError

org.neo4j.graphdb.security.WriteOperationsNotAllowedException

org.neo4j.graphdb.traversal.AlternatingSelectorOrderer

org.neo4j.graphdb.traversal.BidirectionalTraversalDescription

org.neo4j.graphdb.traversal.BidirectionalUniquenessFilter

org.neo4j.graphdb.traversal.BranchCollisionDetector

org.neo4j.graphdb.traversal.BranchCollisionPolicies

org.neo4j.graphdb.traversal.BranchCollisionPolicy

org.neo4j.graphdb.traversal.BranchOrderingPolicies

org.neo4j.graphdb.traversal.BranchOrderingPolicy

org.neo4j.graphdb.traversal.BranchSelector

org.neo4j.graphdb.traversal.BranchState

org.neo4j.graphdb.traversal.Evaluation

org.neo4j.graphdb.traversal.Evaluator.AsPathEvaluator

org.neo4j.graphdb.traversal.Evaluator

org.neo4j.graphdb.traversal.Evaluators

org.neo4j.graphdb.traversal.InitialBranchState.Adapter

org.neo4j.graphdb.traversal.InitialBranchState.State

org.neo4j.graphdb.traversal.InitialBranchState

org.neo4j.graphdb.traversal.LevelSelectorOrderer

org.neo4j.graphdb.traversal.PathEvaluator.Adapter

org.neo4j.graphdb.traversal.PathEvaluator

org.neo4j.graphdb.traversal.Paths.DefaultPathDescriptor

org.neo4j.graphdb.traversal.Paths.PathDescriptor

org.neo4j.graphdb.traversal.Paths

org.neo4j.graphdb.traversal.SideSelectorPolicies

Classes removed or excluded from the public API
org.neo4j.graphdb.traversal.SideSelectorPolicy
org.neo4j.graphdb.traversal.SideSelector
org.neo4j.graphdb.traversal.Sorting
org.neo4j.graphdb.traversal.TraversalBranch
org.neo4j.graphdb.traversal.TraversalContext
org.neo4j.graphdb.traversal.TraversalDescription
org.neo4j.graphdb.traversal.TraversalMetadata
org.neo4j.graphdb.traversal.Traverser
org.neo4j.graphdb.traversal.UniquenessFactory
org.neo4j.graphdb.traversal.UniquenessFilter
org.neo4j.graphdb.traversal.Uniqueness
org.neo4j.helpers.AdvertisedSocketAddress
org.neo4j.helpers.Args.ArgsParser
org.neo4j.helpers.Args.Option
org.neo4j.helpers.Args
org.neo4j.helpers.ArrayUtil.ArrayEquality
org.neo4j.helpers.ArrayUtil
org.neo4j.helpers.Assertion
org.neo4j.helpers.Cancelable
org.neo4j.helpers.CancellationRequest
org.neo4j.helpers.Clock
org.neo4j.helpers.CloneableInPublic
org.neo4j.helpers.Exceptions
org.neo4j.helpers.Format
org.neo4j.helpers.FutureAdapter.Present
org.neo4j.helpers.FutureAdapter
org.neo4j.helpers.HostnamePort
org.neo4j.helpers.ListenSocketAddress
org.neo4j.helpers.Listeners.Notification
org.neo4j.helpers.Listeners
org.neo4j.helpers.MathUtil
org.neo4j.helpers.NamedThreadFactory.Monitor
org.neo4j.helpers.NamedThreadFactory
org.neo4j.helpers.Numbers
org.neo4j.helpers.PortBindException
org.neo4j.helpers.ProcessFailureException.Entry
org.neo4j.helpers.ProcessFailureException
org.neo4j.helpers.Reference
org.neo4j.helpers.RunCarefully
org.neo4j.helpers.Service.Implementation

Classes removed or excluded from the public API

org.neo4j.helpers.Service

org.neo4j.helpers.SocketAddressParser

org.neo4j.helpers.Strings

org.neo4j.helpers.TaskControl

org.neo4j.helpers.TaskCoordinator

org.neo4j.helpers.TextUtil

org.neo4j.helpers.ThisShouldNotHappenError

org.neo4j.helpers.TimeUtil

org.neo4j.helpers.TransactionTemplate.Monitor.Adapter

org.neo4j.helpers.TransactionTemplate.Monitor

org.neo4j.helpers.TransactionTemplate

org.neo4j.helpers.Uris

org.neo4j.helpers.collection.ArrayIterator

org.neo4j.helpers.collection.BoundedIterable

org.neo4j.helpers.collection.CachingIterator

org.neo4j.helpers.collection.CastingIterator

org.neo4j.helpers.collection.CatchingIteratorWrapper

org.neo4j.helpers.collection.CollectorsUtil

org.neo4j.helpers.collection.CombiningIterable

org.neo4j.helpers.collection.CombiningIterator

org.neo4j.helpers.collection.CombiningResourceIterator

org.neo4j.helpers.collection.ExceptionHandlingIterable

org.neo4j.helpers.collection.FilteringIterable

org.neo4j.helpers.collection.FilteringIterator

org.neo4j.helpers.collection.FirstItemIterable

org.neo4j.helpers.collection.IterableWrapper

org.neo4j.helpers.collection.Iterables

org.neo4j.helpers.collection.IteratorWrapper

org.neo4j.helpers.collection.Iterators

org.neo4j.helpers.collection.LimitingResourceIterable

org.neo4j.helpers.collection.LimitingResourceIterator

org.neo4j.helpers.collection.LruCache

org.neo4j.helpers.collection.MapUtil.MapBuilder

org.neo4j.helpers.collection.MapUtil

org.neo4j.helpers.collection.MappingResourceIterator

org.neo4j.helpers.collection.MultiSet

org.neo4j.helpers.collection.NestingIterable

org.neo4j.helpers.collection.NestingIterator

org.neo4j.helpers.collection.NestingResourceIterator

org.neo4j.helpers.collection.PagingIterator

Classes removed or excluded from the public API

org.neo4j.helpers.collection.Pair

org.neo4j.helpers.collection.PrefetchingIterator

org.neo4j.helpers.collection.PrefetchingResourceIterator

org.neo4j.helpers.collection.RangeIterator

org.neo4j.helpers.collection.ResourceClosingIterator

org.neo4j.helpers.collection.ResourceIterableWrapper

org.neo4j.helpers.collection.ReverseArrayIterator

org.neo4j.helpers.collection.Visitable

org.neo4j.helpers.collection.Visitor.SafeGenerics

org.neo4j.helpers.collection.Visitor

org.neo4j.index.lucene.LuceneKernelExtensionFactory.Dependencies

org.neo4j.index.lucene.LuceneKernelExtensionFactory

org.neo4j.index.lucene.LuceneKernelExtension

org.neo4j.index.lucene.LuceneTimeline

org.neo4j.index.lucene.QueryContext

org.neo4j.index.lucene.TimelineIndex

org.neo4j.index.lucene.ValueContext

org.neo4j.index.lucene.unsafe.batchinsert.LuceneBatchInserterIndexProvider

org.neo4j.jmx.Description

org.neo4j.jmx.JmxUtils

org.neo4j.jmx.Kernel

org.neo4j.jmx.ManagementInterface

org.neo4j.jmx.Primitives

org.neo4j.jmx.StoreFile

org.neo4j.jmx.StoreSize

org.neo4j.logging.AbstractLogProvider

org.neo4j.logging.AbstractLog

org.neo4j.logging.AbstractPrintWriterLogger

org.neo4j.logging.BufferingLog

org.neo4j.logging.DuplicatingLogProvider

org.neo4j.logging.DuplicatingLog

org.neo4j.logging.FormattedLog.Builder

org.neo4j.logging.FormattedLogProvider.Builder

org.neo4j.logging.FormattedLogProvider

org.neo4j.logging.FormattedLog

org.neo4j.logging.NullLogProvider

org.neo4j.logging.NullLog

org.neo4j.logging.NullLogger

org.neo4j.logging.PrintStreamLogger

org.neo4j.logging.RotatingFileOutputStreamSupplier.RotationListener

Classes removed or excluded from the public API

org.neo4j.logging.RotatingFileOutputStreamSupplier

org.neo4j.logging.slf4j.Slf4jLogProvider

org.neo4j.logging.slf4j.Slf4jLog

org.neo4j.management.BranchedStoreInfo

org.neo4j.management.BranchedStore

org.neo4j.management.CausalClustering

org.neo4j.management.ClusterDatabaseInfo

org.neo4j.management.ClusterMemberInfo

org.neo4j.management.Diagnostics

org.neo4j.management.HighAvailability

org.neo4j.management.IndexSamplingManager

org.neo4j.management.LockManager

org.neo4j.management.MemoryMapping

org.neo4j.management.Neo4jManager

org.neo4j.management.PageCache

org.neo4j.management.RemoteConnection

org.neo4j.management.TransactionManager

org.neo4j.management.WindowPoolInfo

org.neo4j.procedure.Admin

org.neo4j.procedure.PerformsWrites

org.neo4j.procedure.ProcedureTransaction

org.neo4j.procedure.TerminationGuard

org.neo4j.server.helpers.PropertyTypeDispatcher.PropertyTypeArray

org.neo4j.server.helpers.PropertyTypeDispatcher

org.neo4j.server.plugins.BadPluginInvocationException

org.neo4j.server.plugins.ConfigAdapter

org.neo4j.server.plugins.DefaultPluginManager

org.neo4j.server.plugins.Description

org.neo4j.server.plugins.DisabledPluginManager

org.neo4j.server.plugins.Injectable

org.neo4j.server.plugins.MapTypeCaster

org.neo4j.server.plugins.Name

org.neo4j.server.plugins.ParameterDescriptionConsumer

org.neo4j.server.plugins.ParameterList

org.neo4j.server.plugins.Parameter

org.neo4j.server.plugins.PluginInvocationFailureException

org.neo4j.server.plugins.PluginInvokerProvider

org.neo4j.server.plugins.PluginInvoker

org.neo4j.server.plugins.PluginLifecycle

org.neo4j.server.plugins.PluginLookupException

Classes removed or excluded from the public API

org.neo4j.server.plugins.PluginManager
org.neo4j.server.plugins.PluginPoint
org.neo4j.server.plugins.PluginTarget
org.neo4j.server.plugins.SPIPluginLifecycle
org.neo4j.server.plugins.ServerExtender
org.neo4j.server.plugins.ServerPlugin
org.neo4j.server.plugins.Source
org.neo4j.server.rest.repr.AuthorizationRepresentation
org.neo4j.server.rest.repr.BadInputException
org.neo4j.server.rest.repr.ConstraintDefinitionRepresentation
org.neo4j.server.rest.repr.CypherPlanRepresentation
org.neo4j.server.rest.repr.CypherRepresentationDispatcher
org.neo4j.server.rest.repr.CypherResultRepresentation
org.neo4j.server.rest.repr.CypherStatisticsRepresentation
org.neo4j.server.rest.repr.DatabaseRepresentation
org.neo4j.server.rest.repr.DefaultFormat
org.neo4j.server.rest.repr.DiscoveryRepresentation
org.neo4j.server.rest.repr.EntityRepresentation
org.neo4j.server.rest.repr.ExceptionRepresentation
org.neo4j.server.rest.repr.ExtensionInjector
org.neo4j.server.rest.repr.ExtensionPointRepresentation
org.neo4j.server.rest.repr.FullPath
org.neo4j.server.rest.repr.IndexDefinitionRepresentation
org.neo4j.server.rest.repr.IndexRepresentation
org.neo4j.server.rest.repr.IndexedEntityRepresentation
org.neo4j.server.rest.repr.InputFormatProvider
org.neo4j.server.rest.repr.InputFormat
org.neo4j.server.rest.repr.InvalidArgumentsException
org.neo4j.server.rest.repr.ListRepresentation
org.neo4j.server.rest.repr.ListSerializer
org.neo4j.server.rest.repr.ListWriter
org.neo4j.server.rest.repr.MapRepresentation
org.neo4j.server.rest.repr.MappingRepresentation
org.neo4j.server.rest.repr.MappingSerializer
org.neo4j.server.rest.repr.MappingWriter
org.neo4j.server.rest.repr.MediaTypeNotSupportedException
org.neo4j.server.rest.repr.NodeIndexRepresentation
org.neo4j.server.rest.repr.NodeIndexRootRepresentation
org.neo4j.server.rest.repr.NodeRepresentation
org.neo4j.server.rest.repr.ObjectRepresentation

Classes removed or excluded from the public API

org.neo4j.server.rest.repr.ObjectToRepresentationConverter

org.neo4j.server.rest.repr.OutputFormatProvider

org.neo4j.server.rest.repr.OutputFormat

org.neo4j.server.rest.repr.PathRepresentation

org.neo4j.server.rest.repr.PropertiesRepresentation

org.neo4j.server.rest.repr.RelationshipIndexRepresentation

org.neo4j.server.rest.repr.RelationshipIndexRootRepresentation

org.neo4j.server.rest.repr.RelationshipRepresentation

org.neo4j.server.rest.repr.RepresentationDispatcher

org.neo4j.server.rest.repr.RepresentationExceptionHandlingIterable

org.neo4j.server.rest.repr.RepresentationFormatRepository

org.neo4j.server.rest.repr.RepresentationFormat

org.neo4j.server.rest.repr.RepresentationType

org.neo4j.server.rest.repr.RepresentationWriteHandler

org.neo4j.server.rest.repr.Representation

org.neo4j.server.rest.repr.ScoredEntityRepresentation

org.neo4j.server.rest.repr.ScoredNodeRepresentation

org.neo4j.server.rest.repr.ScoredRelationshipRepresentation

org.neo4j.server.rest.repr.ServerExtensionRepresentation

org.neo4j.server.rest.repr.ServerListRepresentation

org.neo4j.server.rest.repr.StreamingFormat

org.neo4j.server.rest.repr.ValueRepresentation

org.neo4j.server.rest.repr.WeightedPathRepresentation

org.neo4j.server.rest.web.BatchOperationService

org.neo4j.server.rest.web.CollectUserAgentFilter

org.neo4j.server.rest.web.CorsFilter

org.neo4j.server.rest.web.CustomStatusType

org.neo4j.server.rest.web.CypherService

org.neo4j.server.rest.web.DatabaseActions.Provider

org.neo4j.server.rest.web.DatabaseActions.RelationshipDirection

org.neo4j.server.rest.web.DatabaseActions

org.neo4j.server.rest.web.DatabaseMetadataService

org.neo4j.server.rest.web.ExtensionService

org.neo4j.server.rest.web.HttpConnectionInfoFactory

org.neo4j.server.rest.web.InternalJettyServletRequest.RequestData

org.neo4j.server.rest.web.InternalJettyServletRequest

org.neo4j.server.rest.web.InternalJettyServletResponse

org.neo4j.server.rest.web.NoSuchPropertyException

org.neo4j.server.rest.web.NodeNotFoundException

org.neo4j.server.rest.web.PropertyValueException

Classes removed or excluded from the public API

org.neo4j.server.rest.web.RelationshipNotFoundException
org.neo4j.server.rest.web.RestfulGraphDatabase.AmpersandSeparatedCollection
org.neo4j.server.rest.web.RestfulGraphDatabase
org.neo4j.server.rest.web.StreamingBatchOperations
org.neo4j.server.rest.web.Surface
org.neo4j.server.rest.web.TransactionUriScheme
org.neo4j.server.rest.web.TransactionService.TransactionUriBuilder
org.neo4j.server.rest.web.TransactionService
org.neo4j.unsafe.batchinsert.BatchInserterIndexProvider
org.neo4j.unsafe.batchinsert.BatchInserterIndex
org.neo4j.unsafe.batchinsert.BatchInserter
org.neo4j.unsafe.batchinsert.BatchInserters
org.neo4j.unsafe.batchinsert.BatchRelationship

Appendix B: External dependencies

This appendix lists the external dependencies in Neo4j 4.0.

The following table lists the external dependencies in Neo4j 4.0:

Group Id	Artifact Id	Version
com.fasterxml.jackson.core	jackson-annotations	2.10.0
com.fasterxml.jackson.core	jackson-core	2.10.0
com.fasterxml.jackson.core	jackson-databind	2.10.0
com.fasterxml.jackson.jaxrs	jackson-jaxrs-base	2.10.0
com.fasterxml.jackson.jaxrs	jackson-jaxrs-json-provider	2.10.0
com.fasterxml.jackson.module	jackson-module-jaxb-annotations	2.10.0
com.github.ben-manes.caffeine	caffeine	2.8.0
com.github.luben	zstd-jni	1.4.3-1
commons-beanutils	commons-beanutils	1.9.4
commons-collections	commons-collections	3.2.2
commons-configuration	commons-configuration	1.10
commons-io	commons-io	2.6
commons-lang	commons-lang	2.6
commons-logging	commons-logging	1.2
com.profesorfalken	jPowerShell	3.0
com.profesorfalken	WMI4Java	1.6.3
com.sun.activation	jakarta.activation	1.2.1
com.sun.istack	istack-commons-runtime	3.0.8
com.sun.xml.fastinfoset	FastInfoset	1.2.16
com.typesafe.akka	akka-actor_2.12	2.5.22
com.typesafe.akka	akka-cluster_2.12	2.5.22
com.typesafe.akka	akka-cluster-tools_2.12	2.5.22
com.typesafe.akka	akka-coordination_2.12	2.5.22
com.typesafe.akka	akka-distributed-data_2.12	2.5.22
com.typesafe.akka	akka-protobuf_2.12	2.5.22
com.typesafe.akka	akka-remote_2.12	2.5.22
com.typesafe.akka	akka-stream_2.12	2.5.22
com.typesafe	config	1.3.3
com.typesafe	ssl-config-core_2.12	0.3.7
info.picocli	picocli	4.0.4
io.aeron	aeron-client	1.15.1
io.aeron	aeron-driver	1.15.1
io.dropwizard.metrics	metrics-core	4.1.0
io.dropwizard.metrics	metrics-graphite	4.1.0

Group Id	Artifact Id	Version
io.dropwizard.metrics	metrics-jmx	4.1.0
io.netty	netty-all	4.1.35.Final
io.netty	netty	3.10.6.Final
io.projectreactor	reactor-core	3.2.10.RELEASE
io.prometheus	simpleclient_common	0.7.0
io.prometheus	simpleclient_dropwizard	0.7.0
io.prometheus	simpleclient_httpserver	0.7.0
io.prometheus	simpleclient	0.7.0
jakarta.activation	jakarta.activation-api	1.2.1
jakarta.annotation	jakarta.annotation-api	1.3.4
jakarta.ws.rs	jakarta.ws.rs-api	2.1.5
jakarta.xml.bind	jakarta.xml.bind-api	2.3.2
javax.activation	activation	1.1.1
javax.servlet	javax.servlet-api	3.1.0
javax.validation	validation-api	2.0.1.Final
javax.ws.rs	javax.ws.rs-api	2.1.1
javax.xml.bind	jaxb-api	2.3.0
jline	jline	2.14.3
net.java.dev.jna	jna	5.4.0
net.jpountz.lz4	lz4	1.3.0
org.agrona	agrona	0.9.31
org.apache.commons	commons-compress	1.19
org.apache.commons	commons-lang3	3.9
org.apache.commons	commons-text	1.7
org.apache.lucene	lucene-analyzers-common	8.2.0
org.apache.lucene	lucene-codecs	8.2.0
org.apache.lucene	lucene-core	8.2.0
org.apache.lucene	lucene-queryparser	8.2.0
org.apache.shiro	shiro-cache	1.4.1
org.apache.shiro	shiro-config-core	1.4.1
org.apache.shiro	shiro-config-ogdl	1.4.1
org.apache.shiro	shiro-core	1.4.1
org.apache.shiro	shiro-crypto-cipher	1.4.1
org.apache.shiro	shiro-crypto-core	1.4.1
org.apache.shiro	shiro-crypto-hash	1.4.1
org.apache.shiro	shiro-event	1.4.1
org.apache.shiro	shiro-lang	1.4.1
org.bitbucket.inkytonik.kiama	kiama_2.12	2.1.0
org.bouncycastle	bcpkix-jdk15on	1.63
org.bouncycastle	bcprov-jdk15on	1.63

Group Id	Artifact Id	Version
org.eclipse.collections	eclipse-collections-api	10.0.0
org.eclipse.collections	eclipse-collections	10.0.0
org.eclipse.jetty	jetty-client	9.4.17.v20190418
org.eclipse.jetty	jetty-http	9.4.17.v20190418
org.eclipse.jetty	jetty-io	9.4.17.v20190418
org.eclipse.jetty	jetty-security	9.4.17.v20190418
org.eclipse.jetty	jetty-server	9.4.17.v20190418
org.eclipse.jetty	jetty-servlet	9.4.17.v20190418
org.eclipse.jetty	jetty-util	9.4.17.v20190418
org.eclipse.jetty	jetty-webapp	9.4.17.v20190418
org.eclipse.jetty	jetty-xml	9.4.17.v20190418
org.glassfish.hk2.external	jakarta.inject	2.5.0
org.glassfish.hk2	hk2-api	2.5.0
org.glassfish.hk2	hk2-locator	2.5.0
org.glassfish.hk2	hk2-utils	2.5.0
org.glassfish.jaxb	jaxb-runtime	2.3.2
org.glassfish.jaxb	txw2	2.3.2
org.glassfish.jersey.containers	jersey-container-servlet-core	2.29
org.glassfish.jersey.containers	jersey-container-servlet	2.29
org.glassfish.jersey.core	jersey-client	2.29
org.glassfish.jersey.core	jersey-common	2.29
org.glassfish.jersey.core	jersey-server	2.29
org.glassfish.jersey.inject	jersey-hk2	2.29
org.glassfish.jersey.media	jersey-media-jaxb	2.29
org.javassist	javassist	3.22.0-CR2
org.jprocesses	jProcesses	1.6.5
org.jvnet.staxex	stax-ex	1.8.1
org.neo4j.licensing-proxy	zstd-proxy	4.0.0-SNAPSHOT
org.ow2.asm	asm-analysis	7.2
org.ow2.asm	asm	7.2
org.ow2.asm	asm-tree	7.2
org.ow2.asm	asm-util	7.2
org.parboiled	parboiled-core	1.2.0
org.parboiled	parboiled-scala_2.12	1.2.0
org.reactivestreams	reactive-streams	1.0.2
org.rogach	scallop_2.12	2.1.1
org.scala-lang.modules	scala-java8-compat_2.12	0.8.0
org.scala-lang.modules	scala-parser-combinators_2.12	1.1.1
org.scala-lang	scala-library	2.12.7
org.scala-lang	scala-reflect	2.12.7

Group Id	Artifact Id	Version
org.slf4j	slf4j-api	1.7.25
org.slf4j	slf4j-nop	1.7.25